

Qualitative Controller Synthesis for Consumption Markov Decision Processes

No Author Given

No Institute Given

Abstract. Consumption Markov Decision Processes (CMDPs) are probabilistic decision-making models of resource-constrained systems. In a CMDP, the controller possesses a certain amount of a critical resource, such as electric power. Each action of the controller can consume some amount of the resource. Resource replenishment is only possible in special *reload states*, in which the resource level can be reloaded up to the full capacity of the system. The task of the controller is to prevent resource exhaustion, i.e. ensure that the available amount of the resource stays non-negative, while ensuring an additional linear-time property. We study the complexity of strategy synthesis in consumption MDPs with almost-sure Büchi objectives. We show that the problem can be solved in polynomial time. We implement our algorithm, and show that it can efficiently solve CMDPs modelling real-world scenarios.

1 Introduction

In the context of formal methods, controller synthesis typically boils down to computing a strategy in an *agent-environment* model, a nondeterministic state-transition model where some of the nondeterministic choices are resolved by the controller and some by an uncontrollable environment. Such models are typically either two-player graph games with an adversarial environment or Markov decision process (MDPs) with a probabilistic environment; the latter case being apt for modelling statistically predictable environments. In this paper, we consider controller synthesis for *resource-constrained MDPs*, where the computed controller must ensure, in addition to satisfying some linear-time property, that the system's operation is not compromised by a lack of necessary resources.

Resource-Constrained Probabilistic Systems. *Resource-constrained* systems need a supply of some resource (e.g. power) for steady operation: the interruption of the supply can lead to undesirable consequences and has to be avoided. For instance, an autonomous system, e.g. an autonomous electric vehicle (AEV), is not able to draw power directly from an endless source. Instead, it has to rely on an internal storage of the resource, e.g. a battery, which has to be replenished in regular intervals to prevent resource exhaustion. Practical examples of AEVs include driverless cars, drones, or planetary rovers [8]. In these domains, resource failures may cause a costly mission failure and even safety risks. Moreover, operation of autonomous systems is subject to probabilistic uncertainty [57]. Hence, in this paper we study the resource-constrained strategy synthesis problem for MDPs.

Models of Resource-Constrained Systems & Limitations of Current Approaches. There is a substantial body of work on verification of resource-constrained systems [23, 11, 9, 3, 59, 56, 38, 39, 7, 5]. The typical approach is to model them as finite-state systems augmented with an integer-valued counter representing the current *resource level*, i.e. the amount of the resource present in the internal storage. The resource constraint requires that the resource level never drops below zero.¹ In the well-known *energy* model [23, 11], each transition is labelled by an integer, and performing an ℓ -labelled transition results in ℓ being added to the counter. Thus, negative numbers stand for resource consumption while positive ones represent re-charging by the respective amount. Many variants of both MDP and game-based energy models were studied, as detailed in the related work. In particular, [26] considers controller synthesis for energy MDPs with qualitative Büchi and parity objectives. The main limitation of energy-based agent-environment models is that in general, they are not known to admit polynomial-time controller synthesis algorithms. Indeed, already the simplest problem, deciding whether a non-negative energy can be maintained in a two-player energy game, is at least as hard as solving mean-payoff graph games [11]; the complexity of the latter being a well-known open problem [45]. This hardness translates also to MDPs [26], making polynomial-time controller synthesis for energy MDPs impossible without a theoretical breakthrough.

Consumption models, introduced in [14], offer an alternative to energy models. In a consumption model, a non-negative integer, *cap*, represents the maximal amount of the resource the system can hold, e.g. the battery capacity. Each transition is labelled by a non-negative number representing the amount of the resource *consumed* when taking the transition (i.e., taking an ℓ -labelled transition decreases the resource level by ℓ). The resource replenishment is different from the energy approach. The consumption approach relies on the fact that reloads are often *atomic events*, e.g. an AEV plugging into a charging station and waiting to finish the charging cycle. Hence, some states in the consumption model are designated as *reload states*, and whenever the system visits a reload state, the resource level is replenished to the full capacity *cap*. Modelling reloads as atomic events is natural and even advantageous: consumption models typically admit more efficient analysis than energy models [14, 47]. However, consumption models have not yet been considered in the probabilistic setting.

Our Contribution. We study strategy synthesis in consumption MDPs with Büchi objectives. Our main theoretical result is stated in the following theorem.

Theorem 1. *Given a consumption MDP \mathcal{M} with a capacity *cap*, an initial resource level $0 \leq d \leq \text{cap}$, and a set T of accepting states, we can decide, in polynomial time, whether there exists a strategy σ such that when playing according to σ , the following consumption-Büchi objectives are satisfied:*

- *Starting with resource level d , the resource level never² drops below 0.*

¹ In some literature, the level is required to stay positive as opposed to non-negative, but this is only a matter of definition: both approaches are equivalent.

² In our model, this is equivalent to requiring that with probability 1, the resource level never drops below 0.

- With probability 1, the system visits some state in T infinitely often.

Moreover, if such a strategy exists then we can compute, in polynomial time, its polynomial-size representation.

In addition to the theoretical analysis, we implemented the algorithm behind Theorem 1 and evaluated it on several benchmarks, including a model of an AEV navigating the streets of Manhattan, constructed using real-world traffic data. The experiments show that our algorithm is able to efficiently solve large CMDPs, offering a good scalability.

For the sake of clarity, we restrict to proving Theorem 1 for a natural sub-class of MDPs called *decreasing consumption MDPs*, where there are no cycles of zero consumption. The restriction is natural (since in typical resource-constrained systems, each action – even idling – consumes some energy, so zero cycles are unlikely) and greatly simplifies presentation.

Significance. Some comments on Theorem 1 are in order. First, all the numbers in the MDP, and in particular the capacity cap , are encoded in binary. Hence, “polynomial time” means time polynomial in the encoding size of the MDP itself and in $\log(cap)$. In particular, a naive “unfolding” of the MDP, i.e. encoding the resource levels between 0 and cap into the states, does not yield a polynomial-time algorithm, but an exponential-time one, since the unfolded MDP has size proportional to cap . We employ a value-iteration-like algorithm to compute minimal energy levels with which one can achieve the consumption-Büchi objectives.

A similar concern applies to the “polynomial-size representation” of the strategy σ . To satisfy a consumption-Büchi objective, σ generally needs to keep track of the current resource level. Hence, under the standard notion of a finite-memory (FM) strategy (which views FM strategies as transducers), σ would require memory proportional to cap , i.e. a memory exponentially large w.r.t. size of the input. However, we show that for each state s we can partition the integer interval $[0, \dots, cap]$ into polynomially many sub-intervals I_1^s, \dots, I_k^s such that, for each $1 \leq j \leq k$, the strategy σ picks the same action whenever the current state is s and the current resource level is in I_j^s . As such, the end points of the intervals are the only extra knowledge required to represent σ , a representation which we call a *counter selector*. We instrument our main algorithm so as to compute, in polynomial time, a polynomial-size counter selector representing the witness strategy σ .

Finally, we consider linear-time properties encoded by Büchi objectives over the states of the MDP. In essence, we assume that the translation of the specification to the Büchi automaton and its product with the original MDP model of the system were already performed. Probabilistic analysis typically requires the use of deterministic Büchi automata, which cannot express all linear-time properties. However, in this paper we consider qualitative analysis, which can be performed using restricted versions of non-deterministic Büchi automata that are still powerful enough to express all ω -regular languages. Examples of such automata are limit-deterministic Büchi automata [54] or good-for-MDPs automata [41]. Alternatively, consumption MDPs with parity objectives could be reduced to consumption-Büchi MDPs using the standard parity-to-Büchi

MDP construction [25, 33, 32, 30]. We abstract from these aspects and focus on the technical core of our problem, solving consumption-Büchi MDPs.

Consequently, to our best knowledge, we present the first polynomial-time algorithm for controller synthesis in resource-constrained MDPs with ω -regular objectives.

Related Work. There is an enormous body of work on energy models. Stemming from the models introduced in [23, 11], the subsequent work covered energy games with various combinations of objectives [27, 13, 49, 12, 21, 20, 18, 10], energy games with multiple resource types [37, 43, 31, 58, 44, 24, 15, 28] or the variants of the above in the MDP [17, 50], infinite-state [1], or partially observable [34] settings. As argued previously, the controller synthesis within these models is at least as hard as solving mean-payoff games. The paper [29] presents polynomial-time algorithms for non-stochastic energy games with special weight structures. Recently, an abstract algebraic perspective on energy models was presented in [22, 35, 36].

Consumption systems were introduced in [14] in the form of consumption games with multiple resource types. Minimizing mean-payoff in automata with consumption constraints was studied in [16].

Our main result requires, as a technical sub-component, solving the *resource-safety* (or just *safety*) problem in consumption MDPs, i.e. computing a strategy which prevents resource exhaustion. The solution to this problem consists (in principle) of a Turing reduction to the problem of minimum cost reachability in two-player games with non-negative costs. The latter problem was studied in [46], with an extension to arbitrary costs considered in [19] (see also [40]). We present our own, conceptually simple, value-iteration-like algorithm for the problem, which is also used in our implementation.

Elements of resource constrained optimization and minimum-cost reachability are also present in the line of work concerning *energy-utility quantiles* in MDPs [5, 7, 6, 4, 42]. In this setting, there is no reloading in the consumption- or energy-model sense, and the task is typically to minimize the total amount of the resource consumed while maximizing the probability that some other objective is satisfied.

Paper Organization & Outline of Techniques After the preliminaries (Section 2), we present counter selectors in Section 3. The next three sections contain the three main steps of our analysis. In Section 4, we solve the safety problem in consumption MDPs. The technical core of our approach is presented in Section 5, where we solve the problem of *safe positive reachability*: finding a resource-safe strategy which ensures that the set T of accepting states is visited with positive probability. Solving consumption-Büchi MDPs then, in principle, consists of repeatedly applying a strategy for safe positive reachability of T , ensuring that the strategy is “re-started” whenever the attempt to reach T fails. Details are given in Section 6. Finally, Section 7 presents our experiments. Due to space constraints, most technical proofs were moved to the appendix.

2 Preliminaries

We denote by \mathbb{N} the set of all non-negative integers and by $\overline{\mathbb{N}}$ the set $\mathbb{N} \cup \{\infty\}$. Given a set I and a vector $\mathbf{v} \in \overline{\mathbb{N}}^I$ of integers indexed by I , we use $\mathbf{v}(i)$ to denote the i -component

of \mathbf{v} . We assume familiarity with basic notions of probability theory. In particular, a *probability distribution* on an at most countable set X is a function $f: X \rightarrow [0, 1]$ s.t. $\sum_{x \in X} f(x) = 1$. We use $\mathcal{D}(X)$ to denote the set of all probability distributions on X .

Definition 1 (CMDP). A consumption Markov decision process (CMDP) is a tuple $\mathcal{M} = (S, A, \Delta, C, R, \text{cap})$ where S is a finite set of states, A is a finite set of actions, $\Delta: S \times A \rightarrow \mathcal{D}(S)$ is a total transition function, $C: S \times A \rightarrow \mathbb{N}$ is a total consumption function, $R \subseteq S$ is a set of reload states where the resource can be reloaded, and cap is a resource capacity.

Given a set $R' \subseteq S$, we denote by $\mathcal{M}(R')$ the CMDP obtained from \mathcal{M} by changing the set of reloads to R' . Given $s \in S$ and $a \in A$, we denote by $\text{Succ}(s, a)$ the set $\{t \mid \Delta(s, a)(t) > 0\}$. A *path* is a (finite or infinite) state-action sequence $\alpha = s_1 a_1 s_2 a_2 s_3 \dots \in (S \times A)^\omega \cup (S \times A)^* \cdot S$ such that $s_{i+1} \in \text{Succ}(s_i, a_i)$ for all i . We define $\alpha_i = s_i$ and $\text{Act}^i(\alpha) = a_i$. We use $\alpha_{..i}$ for the finite prefix $s_1 a_1 s_2 \dots s_i$ of α , we use $\alpha_{i..}$ for the suffix $s_i a_i s_{i+1} \dots$, and $\alpha_{i..j}$ for the infix $s_i a_i \dots s_j$. The *length* of a path α is the number $\text{len}(\alpha)$ of actions on α ($\text{len}(\alpha) = \infty$ if α is infinite).

A finite path α is *simple* if no state appears more than once on α . A finite path is a *cycle* if it starts and ends in the same state. A CMDP is *decreasing* if every simple cycle $s_1 a_1 s_2 \dots a_{k-1} s_k$ there exists $1 \leq i < k$ such that $C(s_i, a_i) > 0$. Throughout this paper we consider only decreasing CMDPs. The only place where this assumption is used are the proofs of Theorem 4 and Theorem 8.

An infinite path is called a *run*. We typically name runs by variants of the symbol ϱ . The set of all runs in \mathcal{M} is denoted $\text{Runs}_{\mathcal{M}}$ or simply Runs if \mathcal{M} is clear from context. A finite path is also called *history*. The set of all possible histories of \mathcal{M} is $\text{hist}_{\mathcal{M}}$ or simply hist . We denote by $\text{last}(\alpha)$ the last state of a history α . Let α be a history with $\text{last}(\alpha) = s_1$ and $\beta = s_1 a_1 s_2 a_2 \dots$; we define a *joint path* as $\alpha \odot \beta = \alpha a_1 s_2 a_2 \dots$.

A *strategy* for \mathcal{M} is a function $\sigma: \text{hist}_{\mathcal{M}} \rightarrow A$ assigning to each history an action to play. A strategy is *memoryless* if $\sigma(\alpha) = \sigma(\beta)$ whenever $\text{last}(\alpha) = \text{last}(\beta)$, i.e., when the decision depends only on the current state. We do not consider randomized strategies in this paper, as they are non-necessary for qualitative ω -regular objectives on finite MDPs [33, 32, 30].

A computation of \mathcal{M} under the control of a given strategy σ from some initial state $s \in S$ creates a path. The path starts with $s_1 = s$. Assume that the current path is α and let $s_i = \text{last}(\alpha)$ (we say that \mathcal{M} is currently in s_i). Then the next action on the path is $a_i = \sigma(\alpha)$ and the next state s_{i+1} is chosen randomly according to $\Delta(s_i, a_i)$. Repeating this process *ad infinitum* yields an infinite sample run ϱ . We say that a ϱ is σ -compatible if it can be produced using this process, and s -initiated if it starts in s . We denote the set of all σ -compatible s -initiated runs by $\text{Comp}_{\mathcal{M}}(\sigma, s)$.

We denote by $\mathbb{P}_{\mathcal{M}, s}^\sigma(A)$ the probability that a sample run from $\text{Comp}_{\mathcal{M}}(\sigma, s)$ belongs to a given measurable set of runs A (the subscript \mathcal{M} is dropped when \mathcal{M} is known from the context). For details on the formal construction of measurable sets of runs as well as the probability measure $\mathbb{P}_{\mathcal{M}, s}^\sigma$ see [2].

2.1 Resource: Consumption, Levels, and Objectives

We denote by $\text{cap}(\mathcal{M})$ the battery capacity in the MDP \mathcal{M} . A resource is consumed along paths and can be reloaded in the reload states up to the full capacity. For a path $\alpha = s_1 a_1 s_2 \dots$ we define the consumption of α as $\text{cons}(\alpha) = \sum_{i=1}^{\text{len}(\alpha)} C(s_i, a_i)$ (since the consumption is non-negative, the sum is always well defined, though possibly diverging). Note that cons does not consider reload states at all. To accurately track the remaining amount of the resource, we use the concept of a *resource level*.

Definition 2 (Resource level). Let \mathcal{M} be a CMDP with a set of reload states R , let α be a history, and let $0 \leq d \leq \text{cap}(\mathcal{M})$ be an integer called initial load. Then the energy level after α initialized by d , denoted by $RL_d^{\mathcal{M}}(\alpha)$ or simply as $RL_d(\alpha)$, is defined inductively as follows: for a zero-length history s we have $RL_d^{\mathcal{M}}(s) = d$. For a non-zero-length history $\alpha = \beta a$ we denote $c = C(\text{last}(\beta), a)$, and put

$$RL_d^{\mathcal{M}}(\alpha) = \begin{cases} RL_d^{\mathcal{M}}(\beta) - c & \text{if } \text{last}(\beta) \notin R \text{ and } c \leq RL_d^{\mathcal{M}}(\beta) \neq \perp \\ \text{cap}(\mathcal{M}) - c & \text{if } \text{last}(\beta) \in R \text{ and } c \leq \text{cap}(\mathcal{M}) \text{ and } RL_d^{\mathcal{M}}(\beta) \neq \perp \\ \perp & \text{otherwise} \end{cases}$$

Let α be a history and let $f, l \geq 0$ that are the minimal and maximal indices i such that $\alpha_i \in R$, respectively. Following the inductive definition of $RL_d(\alpha)$ it is easy to see that if we have $RL_d(\alpha) \neq \perp$, then $RL_d(\alpha_{..i}) = d - \text{cons}(\alpha_{..i})$ holds for all $i \leq f$ and $RL_d(\alpha) = \text{cap}(\mathcal{M}) - \text{cons}(\alpha_{l..})$. Further, for each history α and d such that $e = RL_d(\alpha) \neq \perp$, and each history β suitable for joining with α it holds that $RL_d(\alpha \odot \beta) = RL_e(\beta)$.

A run ϱ is *d-safe* if and only if the energy level initialized by d is a non-negative number for each finite prefix of ϱ , i.e. if for all $i > 0$ we have $RL_d(\varrho_{..i}) \neq \perp$. We say that a run is safe if it is $\text{cap}(\mathcal{M})$ -safe. The next lemma follows immediately from the definition of an energy level.

Lemma 1. Let $\varrho = s_1 a_1 s_2 \dots$ be a *d-safe* run for some d and let α be a history such that $\text{last}(\alpha) = s_1$. Then the run $\alpha \odot \varrho$ is *e-safe* if $RL_e(\alpha) \geq d$.

Objectives AN objective is a set of runs. The objective $\text{SafeRuns}(d)$ contains exactly *d-safe* runs. Given a target set $T \subseteq S$ and $i \in \mathbb{N}$, we define $\text{Reach}_T^i = \{\varrho \in \text{Runs} \mid \varrho_j \in T \text{ for some } 1 \leq j \leq i+1\}$ to be the set of all runs that reach some state from T within the first i steps. We put $\text{Reach}_T = \bigcup_{i \in \mathbb{N}} \text{Reach}_T^i$. Finally, the set $\text{Büchi}_T = \{\varrho \in \text{Runs} \mid \varrho_i \in T \text{ for infinitely many } i \in \mathbb{N}\}$.

Problems We solve three main qualitative problems for CMDPs, namely *safety*, *positive reachability*, and *Büchi*.

Let us fix a state s and a target set of states T . We say that a strategy is *d-safe* in s if $\text{Comp}(\sigma, s) \subseteq \text{SafeRuns}(d)$. We say that σ is *T-positive d-safe* in s if it is *d-safe* in s and $\mathbb{P}_s^\sigma(\text{Reach}_T) > 0$, which means that there exists a run in $\text{Comp}(\sigma, s)$ that visits T . Finally, we say that σ is *T-Büchi d-safe* in a state s if it is *d-safe* in s and $\mathbb{P}_s^\sigma(\text{Büchi}_T) = 1$.

The vectors *Safe*, *SafePR_T* (PR for “positive reachability”), and *SafeBüchi_T* of type $\overline{\mathbb{N}}^S$ contain, for each $s \in S$, the minimal d such that there exists a strategy that is *d-safe*

in s , T -positive d -safe in s , and T -Büchi d -safe in s , respectively, and ∞ if no such strategy exists.

The problems we consider for a given CMDP are:

- *Safety*: compute the vector Safe and a strategy that is $\text{Safe}(s)$ -safe in every $s \in S$.
- *Positive reachability*: compute the vector SafePR_T and a strategy that is T -positive $\text{SafePR}_T(s)$ -safe in every state s .
- *Büchi*: compute SafeBüchi_T and a strategy that is T -Büchi $\text{SafeBüchi}_T(s)$ -safe in every state s .




Fig. 1. CMDP $\mathcal{M} = (\{s_1, s_2, s_3, s_4, s_5\}, \{a_1, a_2\}, \Delta, C, \{s_2, s_5\}, 20)$ where distributions in Δ are indicated by gray numbers (we leave out 1 when an action has only one successor), and the cost of an action follows its name in the edge labels. Actions labeled by a_i represent that Δ and C are defined identically for both actions a_1 and a_2 . The blue background indicates a target set $T = \{s_2\}$, while the double circles represent the reload states.

We illustrate the key concepts using the example CMDP \mathcal{M} in Figure 1. Consider the parameterized history $\alpha^i = (s_1 a_2 s_5 a_2)^i s_1$. Then $\text{cons}(\alpha^i) = 3i$ while $RL_2(\alpha^i) = 19$ for all $i \geq 1$. Thus, a strategy, that always picks a_2 in s_1 is d -safe in s_1 for all $d \geq 2$. On the other hand, a strategy that always picks a_1 in s_1 is *not* d -safe in s_1 for any $0 \leq d \leq 20$. Now consider again the strategy that always picks a_2 ; such a strategy is 2-safe in s_1 , but is not useful if we attempt to eventually reach T . Hence memoryless strategies are not sufficient in our setting. Consider instead a strategy σ that, in s_1 , picks a_1 whenever the current resource level is at least 10 and picks a_2 otherwise. Such a strategy is 2-safe in s_1 and guarantees reaching s_2 with a positive probability: we need at least 10 units of energy to return to s_5 in the case we are unlucky and picking a_1 leads us to s_3 . If we are lucky, a_1 leads us to s_2 by consuming just 5 units of the resource, witnessing that σ is T -positive. As a matter of fact, during *every* revisit of s_5 there is a $\frac{1}{2}$ chance of hitting s_2 during the next try, so σ actually ensures that s_2 is visited with probability 1.

We note that solving a CMDP is very different from solving a consumption 2-player game [14]. Indeed, imagine that in Figure 1, the outcome of the action a_1 from state s_1 is resolved by an adversarial player. In such a game, there is no strategy that would guarantee reaching T at all.

The strategy σ we discussed above uses finite memory to track the resource level exactly. An efficient representation of such strategies is described in the next section.

3 Counter Strategies

In this section, we define a succinct representation of finite-memory strategies via so called counter selectors. Under the standard definition, a strategy σ is a *finite memory* strategy, if σ can be encoded by a *memory structure*, a type of finite transducer. Formally, a memory structure is a tuple $\mu = (M, \text{nxt}, \text{up}, m_0)$ where M is a finite set of *memory elements*, $\text{nxt}: M \times S \rightarrow A$ is a *next action* function, $\text{up}: M \times S \times A \times S \rightarrow M$ is a *memory update* function, and $m_0: S \rightarrow M$ is the *memory initialization function*. The function up can be lifted to a function $\text{up}^*: M \times \text{hist} \rightarrow M$ as follows.

$$\text{up}^*(m, \alpha) = \begin{cases} m & \text{if } \alpha = s \text{ has length } 0 \\ \text{up}(\text{up}^*(m, \beta), \text{last}(\beta), a, t) & \text{if } \alpha = \beta a t \text{ for some } a \in A \text{ and } t \in S \end{cases}$$

The structure μ encodes a strategy σ_μ such that for each history $\alpha = s_1 a_1 s_2 \dots s_n$ we have $\sigma_\mu(\alpha) = \text{nxt}(\text{up}^*(m_0(s_1), \alpha), s_n)$.

In our setting, strategies need to track energy levels of histories. Let us fix an CMDP $\mathcal{M} = (S, A, \Delta, C, R, \text{cap})$. A non-exhausted energy level is always a number between 0 and $\text{cap}(\mathcal{M})$, which can be represented with a binary-encoded bounded counter. We call strategies with such counters *finite counter (FC) strategies*. An FC strategy selects actions to play according to *selection rules*.

Definition 3 (Selection rule). A selection rule φ for \mathcal{M} is a partial function from the set $\{0, \dots, \text{cap}(\mathcal{M})\}$ to A . Undefined value for some n is indicated by $\varphi(n) = \perp$.

We use $\text{dom}(\varphi) = \{n \in \{0, \dots, \text{cap}(\mathcal{M})\} \mid \varphi(n) \neq \perp\}$ to denote the domain of φ and we use $\text{Rules}_{\mathcal{M}}$ or simply Rules for the set of all selection rules for \mathcal{M} . Intuitively, a selection according to rule φ selects the action that corresponds to the largest value from $\text{dom}(\varphi)$ that is not larger than the current energy level. To be more precise, if $\text{dom}(\varphi)$ consists of numbers $n_1 < n_2 < \dots < n_k$, then the action to be selected in a given moment is $\varphi(n_i)$, where n_i is the largest element of $\text{dom}(\varphi)$ which is less than or equal to the current amount of the resource. In other words, $\varphi(n_i)$ is to be selected if the current resource level is in $[n_i, n_{i+1})$ (putting $n_{k+1} = \infty$).

Definition 4 (Counter selector). A counter selector for \mathcal{M} is a function $\Sigma: S \rightarrow \text{Rules}$.

A counter selector itself is not enough to describe a strategy. A strategy needs to keep track of the energy level throughout the path. With a vector $\mathbf{r} \in \{0, \dots, \text{cap}(\mathcal{M})\}^S$ of initial resource levels, each counter selector Σ defines a strategy $\Sigma^{\mathbf{r}}$ that is encoded by the following memory structure $(M, \text{nxt}, \text{up}, m_0)$ with $a \in A$ being a globally fixed action (for uniqueness). We stipulate that $\perp < n$ for all $n \in \mathbb{N}$.

- $M = \{\perp\} \cup \{0, \dots, \text{cap}(\mathcal{M})\}$.
- Let $m \in M$ be a memory element, let $s \in S$ be a state, let $n \in \text{dom}(\Sigma(s))$ be the largest element of $\text{dom}(\Sigma(s))$ such that $n \leq m$. Then $\text{nxt}(m, s) = \Sigma(s)(n)$ if n exists, and $\text{nxt} = a$ otherwise.
- The function up is defined for each $m \in M, a \in A, s, t \in S$ as follows.

$$\text{up}(m, s, a, t) = \begin{cases} m - C(s, a) & \text{if } s \notin R \text{ and } C(s, a) \leq m \neq \perp \\ \text{cap}(\mathcal{M}) - C(s, a) & \text{if } s \in R \text{ and } C(s, a) \leq \text{cap}(\mathcal{M}) \text{ and } m \neq \perp \\ \perp & \text{otherwise.} \end{cases}$$

- The function m_0 is $m_0(s) = \mathbf{r}(s)$.

A strategy σ is a finite counter (FC) strategy if there is a counter selector Σ and a vector \mathbf{r} such that $\sigma = \Sigma^{\mathbf{r}}$. The counter selector can be imagined as a finite-state device that implements σ using $O(\log(\text{cap}(\mathcal{M})))$ bits of additional memory (counter) used to represent numbers $0, 1, \dots, \text{cap}(\mathcal{M})$. The device uses the counter to keep track of the current resource level, the element \perp representing energy exhaustion. Note that a counter selector can be exponentially more succinct than the corresponding memory structure.

4 Safety

In this section, we present an algorithm that computes, for each state, the minimal value d (if it exists) such that there exists a d -safe strategy from that state. We also provide the corresponding strategy. In the remainder of the section we fix an MDP \mathcal{M} .

A d -safe run has the following two properties: (i) It consumes at most d units of the resource (energy) before it reaches the first reload state, and (ii) it never consumes more than $\text{cap}(\mathcal{M})$ units of the resource between 2 visits of reload states. To ensure (ii), we need to identify a maximal subset $R' \subseteq R$ of reload states for which there is a strategy σ that, starting in some $r \in R'$, can always reach R' again (within at least one step) using at most $\text{cap}(\mathcal{M})$ resource units. The d -safe strategy we seek can be then assembled from σ and from a strategy that suitably navigates towards R' , which is needed for (i).

In the core of both properties (i) and (ii) lies the problem of *minimum cost reachability*. Hence, in the next subsection, we start with presenting necessary results on this problem.

4.1 Minimum Cost Reachability

The problem of minimum cost reachability with non-negative costs was studied before [46]. Here we present a simple approach to the problem used in our implementation. Hence, most of the technical details are deferred to the appendix.

Definition 5. Let $T \subseteq S$ be a set of target states, let $\alpha = s_1 a_1 s_2 \dots$ be a finite or infinite path, and let $1 \leq f$ be the smallest index such that $s_f \in T$. We define consumption of α to T as $\text{ReachCons}_{\mathcal{M},T}(\alpha) = \text{cons}(\alpha_{..f})$ if f exists and we set $\text{ReachCons}_{\mathcal{M},T}(\alpha) = \infty$ otherwise. For a strategy σ and a state $s \in S$ we define $\text{ReachCons}_{\mathcal{M},T}(\sigma, s) = \sup_{\varrho \in \text{Comp}(\sigma, s)} \text{ReachCons}_{\mathcal{M},T}(\varrho)$.

A minimum cost reachability of T from s is a vector defined as

$$\text{MinReach}_{\mathcal{M},T}(s) = \inf \{ \text{ReachCons}_{\mathcal{M},T}(\sigma, s) \mid \sigma \text{ is a strategy for } \mathcal{M} \}.$$

As usual, we drop the subscript \mathcal{M} when \mathcal{M} is clear from context. Intuitively, $d = \text{MinReach}_T(s)$ is the minimal initial load with which some strategy can ensure reaching T with consumption at most d , when starting in s . We say that a strategy σ is optimal for MinReach_T if we have that $\text{MinReach}_T(s) = \text{ReachCons}_T(\sigma, s)$ for all states $s \in S$.

We also define functions $\text{ReachCons}_{\mathcal{M},T}^+$ and the vector $\text{MinReach}_{\mathcal{M},T}^+$ in a similar fashion with one exception: we require the index f from definition of $\text{ReachCons}_{\mathcal{M},T}(\alpha)$ to be strictly larger than 1, which enforces to take at least one step to reach T .

For the rest of this section, fix a target set T and consider the following functional \mathcal{F} :

$$\mathcal{F}(\mathbf{v})(s) = \begin{cases} \min_{a \in A} (C(s, a) + \max_{t \in \text{Succ}(s, a)} \mathbf{v}(t)) & s \notin T \\ 0 & s \in T \end{cases}$$

\mathcal{F} is a simple generalization of the standard Bellman functional used for computing shortest paths in graphs. The proof of the following Theorem is rather standard and is deferred to Appendix A.1.

Theorem 2. *Denote by n the length of the longest simple path in \mathcal{M} . Let \mathbf{x}_T be a vector such that $\mathbf{x}_T(s) = 0$ if $s \in T$ and $\mathbf{x}_T(s) = \infty$ otherwise. Then iterating \mathcal{F} on \mathbf{x}_T yields a fixpoint in at most n steps and this fixpoint equals MinReach_T .*

To compute $\text{MinReach}_{\mathcal{M}, T}^+$, we construct a new CMDP $\widetilde{\mathcal{M}}$ from \mathcal{M} by adding a copy \tilde{s} of each state $s \in S$ such that dynamics in \tilde{s} is the same as in s ; i.e. for each $a \in A$, $\Delta(\tilde{s}, a) = \Delta(s, a)$ and $C(\tilde{s}, a) = C(s, a)$. We denote the new state set as \widetilde{S} . We don't change the set of reload states, so \tilde{s} is *never* in T , even if s is. The following lemma is straightforward.

Lemma 2. *Let \mathcal{M} be a CMDP and let $\widetilde{\mathcal{M}}$ be the CMDP constructed as above. Then for each state s of \mathcal{M} it holds $\text{MinReach}_{\mathcal{M}, T}^+(s) = \text{MinReach}_{\widetilde{\mathcal{M}}, T}(\tilde{s})$.*

4.2 Safely Reaching Reload States

In the following, we use $\text{MinInitCons}_{\mathcal{M}}$ (read *minimal initial consumption*) for the vector $\text{MinReach}_{\mathcal{M}, R}^+$ – minimal resource level that ensures we can surely reach a reload state in at least one step. By Lemma 2 and Theorem 2 we can construct $\widetilde{\mathcal{M}}$ and iterate the operator \mathcal{F} for $|S|$ steps to compute $\text{MinInitCons}_{\mathcal{M}}$. Note that S is the state space of \mathcal{M} since introducing the new states into $\widetilde{\mathcal{M}}$ did not increase the length of the maximal simple path. However, we can avoid the construction of $\widetilde{\mathcal{M}}$ and still compute $\text{MinInitCons}_{\mathcal{M}}$ using a *truncated* version of the functional \mathcal{F} , which is the approach used in our implementation. We first introduce the following truncation operator:

$$\|\mathbf{x}\|_{\mathcal{M}}(s) = \begin{cases} \mathbf{x}(s) & \text{if } s \notin R, \\ 0 & \text{if } s \in R. \end{cases}$$

Then, we define a truncated functional \mathcal{G} as follows:

$$\mathcal{G}(\mathbf{v})(s) = \min_{a \in A} \left(C(s, a) + \max_{s' \in \text{Succ}(s, a)} \|\mathbf{v}\|_{\mathcal{M}}(s') \right).$$

The following lemma connects the iteration of \mathcal{G} on \mathcal{M} with the iteration of \mathcal{F} on $\widetilde{\mathcal{M}}$.

Lemma 3. *Let $\infty \in \overline{\mathbb{N}}^S$ be a vectors with all components equal to ∞ . Consider iterating \mathcal{G} on ∞ in \mathcal{M} and \mathcal{F} on \mathbf{x}_R in $\widetilde{\mathcal{M}}$. Then for each $i \geq 0$ and each $s \in R$ we have $\mathcal{G}^i(\infty)(s) = \mathcal{F}^i(\mathbf{x}_R)(\tilde{s})$ and for every $s \in S \setminus R$ we have $\mathcal{G}^i(\infty)(s) = \mathcal{F}^i(\mathbf{x}_R)(s)$.*

Algorithm 1: Algorithm for computing $MinInitCons_{\mathcal{M}}$.

Input: CMDP $\mathcal{M} = (S, A, \mathcal{A}, C, R, cap)$
Output: The vector $MinInitCons_{\mathcal{M}}$

```

1 initialize  $\mathbf{x} \in \overline{\mathbb{N}}^S$  to be  $\infty$  in every component;
2 repeat
3    $\mathbf{x}_{old} \leftarrow \mathbf{x}$ ;
4   foreach  $s \in S$  do
5      $c \leftarrow \min_{a \in A} \{C(s, a) + \max_{s' \in Succ(s, a)} \|\mathbf{x}_{old}\|_{\mathcal{M}}(s')\}$ ;
6     if  $c < \mathbf{x}(s)$  then
7        $\mathbf{x}(s) \leftarrow c$ ;
8 until  $\mathbf{x}_{old} = \mathbf{x}$ ;
9 return  $\mathbf{x}$ 

```

Algorithm 1 uses \mathcal{G} to compute the vector $MinInitCons_{\mathcal{M}}$.

Theorem 3. *Algorithm 1 correctly computes the vector $MinInitCons_{\mathcal{M}}$. Moreover, the repeat-loop terminates after at most $|S|$ iterations.*

Proof. The repeat-loop performs the iteration of the operator \mathcal{G} . We show that the fixed point of the iteration equals $MinInitCons_{\mathcal{M}}$. Consider the iteration of \mathcal{F} and \mathcal{G} on ∞ and \mathbf{x}_R , respectively. Let i be the number of steps (possibly infinite) after which the \mathcal{F} -iteration reaches a fixed point and j the number of steps after which the \mathcal{G} -iteration reaches a fixed point. We prove that $i = j$. Indeed, for each step k we have that

$$\mathcal{G}^{k+1}(\infty)(s) \neq \mathcal{G}^k(\infty)(s) \Leftrightarrow \mathcal{F}^{k+1}(\mathbf{x}_R)(\tilde{s}) \neq \mathcal{F}^k(\mathbf{x}_R)(\tilde{s}) \quad (1)$$

(by Lemma 3). Hence, $i \geq j$. For the reverse inequality, assume that $i > j$. Then there is $t \in \tilde{S}$ such that $(\mathbf{x}_R)(t) \neq \mathcal{F}^j(\mathbf{x}_R)(t)$. From (1) and from the fact that the \mathcal{G} -iteration already reached a fixed point we get that $t \in S$. Then either $t \in S \setminus R$, but then by Lemma 3 we have $\mathcal{G}^{j+1}(\infty)(t) = \mathcal{F}^{j+1}(\mathbf{x}_R)(t) \neq \mathcal{F}^j(\mathbf{x}_R)(t) = \mathcal{G}^j(\infty)(t)$ a contradiction with \mathcal{G} -iteration already being at a fixed point. Or $t \in S \cap R$, but then $\mathcal{F}^{j+1}(\mathbf{x}_R)(t) = \mathcal{F}^j(\mathbf{x}_R)(t) = 0$, again a contradiction.

Hence, iterating \mathcal{G} also reaches a fixed point in at most $|S|$ -steps, by Theorem 2. Moreover, for each $s \in S$ we have $\mathcal{G}^i(\infty)(s) = \mathcal{F}^i(\mathbf{x}_R)(\tilde{s}) = MinReach_{\tilde{\mathcal{M}}, R}(\tilde{s}) = MinInitCons_{\mathcal{M}}(s)$, the first equality coming from Lemma 3, the second from Theorem 2 and from the fact that $i = j$ and the last one from Lemma 2. \square

4.3 Solving the Safety Problem

We want to identify a set $R' \subseteq R$ such that we can reach R' in at least 1 step and with consumption at most $cap = cap(\mathcal{M})$, from each $r \in R'$. This entails identifying the maximal $R' \subseteq R$ such that $MinInitCons_{\mathcal{M}(R')} \leq cap$ for each $r \in R'$. This can be done by initially setting $R' = R$ and iteratively removing states that have $MinInitCons_{\mathcal{M}(R')} > cap$, from R' , as in Algorithm 2.

Algorithm 2: Computing the vector $\text{Safe}_{\mathcal{M}}$.

Input: CMDP \mathcal{M}
Output: The vector $\text{Safe}_{\mathcal{M}}$

```
1  $cap \leftarrow cap(\mathcal{M});$ 
2  $Rel \leftarrow R; ToRemove \leftarrow \emptyset;$ 
3 repeat
4    $Rel \leftarrow Rel \setminus ToRemove;$ 
5    $\mathbf{mic} \leftarrow \text{MinInitCons}_{\mathcal{M}(Rel)};$ 
6    $ToRemove \leftarrow \{r \in Rel \mid \mathbf{mic}(r) > cap\};$ 
7 until  $ToRemove = \emptyset;$ 
8 foreach  $s \in S$  do
9   if  $\mathbf{mic}(s) > cap$  then  $\mathbf{out}(s) = \infty;$ 
10  else  $\mathbf{out}(s) = \mathbf{mic}(s);$ 
11 return out
```

Theorem 4. Algorithm 2 computes the vector $\text{Safe}_{\mathcal{M}}$ in polynomial time.

Proof. The algorithm clearly terminates. Computing $\text{MinInitCons}_{\mathcal{M}(Rel)}$ on line 5 takes a polynomial number of steps per call due to Theorem 3 and since $\mathcal{M}(Rel)$ has asymptotically the same size as \mathcal{M} . Since the repeat loop performs at most $|R|$ iterations, the complexity follows.

As for correctness, we first prove that $\mathbf{out} \leq \text{Safe}_{\mathcal{M}}$. It suffices to prove for each $s \in S$ that upon termination, $\mathbf{mic}(s) \leq \text{Safe}_{\mathcal{M}}(s)$ whenever the latter value is finite. Since $\text{MinInitCons}_{\mathcal{M}'}(s) \leq \text{Safe}_{\mathcal{M}'}(s)$ for each MDP \mathcal{M}' and each its state such that $\text{Safe}_{\mathcal{M}'}(s) < \infty$, it suffices to show that $\text{Safe}_{\mathcal{M}(Rel)} \leq \text{Safe}_{\mathcal{M}}$ is an invariant of the algorithm (as a matter of fact, we prove that $\text{Safe}_{\mathcal{M}(Rel)} = \text{Safe}_{\mathcal{M}}$). To this end, it suffices to show that at every point of execution $\text{Safe}_{\mathcal{M}}(t) = \infty$ for each $t \in R \setminus Rel$: indeed, if this holds, no strategy that is safe for some state $s \neq t$ can play an action a from s such that $t \in \text{Succ}(s, a)$, so declaring such states non-reloading does not influence the $\text{Safe}_{\mathcal{M}}$ -values. So denote by Rel_i the contents of Rel after the i -th iteration. We prove, by induction on i , that $\text{Safe}_{\mathcal{M}}(s) = \infty$ for all $s \in R \setminus Rel_i$. For $i = 0$ we have $R = Rel$, so the statement holds. For $i > 0$, let $s \in R \setminus Rel_i$, and let σ be any strategy. If some run from $\text{Comp}(\sigma, s)$ visits a state from $R \setminus Rel_{i-1}$, then σ is not cap -safe, by induction hypothesis. Now assume that all such runs only visit reload states from Rel_{i-1} . Then, since $\text{MinInitCons}_{\mathcal{M}(Rel_{i-1})}(s) > cap$, there must be a run $\varrho \in \text{Comp}(\sigma, s)$ with $\text{ReachCons}_{Rel_{i-1}}^+(\varrho) > cap$. Assume that ϱ is cap -safe in s . Since we consider only decreasing CMDPs, ϱ must infinitely often visit a reload state (as it cannot get stuck in a zero cycle). Hence, there exists an index $f > 1$ such that $\varrho_f \in Rel_{i-1}$, and for this f we have $RL_{cap}(\varrho..f) = \perp$, a contradiction. So again, σ is not safe in s . Since there is no safe strategy from s , we have $\text{Safe}_{\mathcal{M}}(s) = \infty$.

Finally, we need to prove that upon termination, $\mathbf{out} \geq \text{Safe}_{\mathcal{M}}$. Informally, per the definition of \mathbf{out} , from every state s we can ensure reaching a state of Rel by consuming at most $\mathbf{out}(s)$ units of the resource. Once in Rel , we can ensure that we can again return to Rel without consuming more than cap units of the resource. Hence, when

starting with $\mathbf{out}(s)$ units, we can surely prevent resource exhaustion. The details are in Appendix A.3. \square

Definition 6. We call an action a safe in a state s if one of the following conditions holds:

- $s \notin R$ and $C(s, a) + \max_{t \in \text{Succ}(s, a)} \text{Safe}_{\mathcal{M}}(t) \leq \text{Safe}_{\mathcal{M}}(s)$; or
- $s \in R$ and $C(s, a) + \max_{t \in \text{Succ}(s, a)} \text{Safe}_{\mathcal{M}}(t) \leq \text{cap}(\mathcal{M})$.

Note that by the definition of $\text{Safe}_{\mathcal{M}}$, for each state s with $\text{Safe}_{\mathcal{M}}(s) < \infty$ there is always at least one action safe in s . For states s s.t. $\text{Safe}_{\mathcal{M}}(s) = \infty$, we stipulate all actions to be safe in s .

Theorem 5. Any strategy which always selects an action that is safe in the current state is $\text{Safe}_{\mathcal{M}}(s)$ -safe in every state s . In particular, in each consumption MDP \mathcal{M} there is a memoryless strategy σ that is $\text{Safe}_{\mathcal{M}}(s)$ -safe in every state s . Moreover, σ can be computed in polynomial time.

Proof. The first part of the theorem follows directly from Definition 6, Definition 2 (resource levels), and from definition of d -safe runs. The second part is a corollary of Theorem 4 and the fact that in each state, the safe strategy from Definition 6 can fix one such action in each state and thus is memoryless. The complexity follows from Theorem 4. \square

5 Positive Reachability

In this section, we focus on strategies that are safe and such that at least one run they produce visits a given set $T \subseteq S$ of *targets*. The main contribution of this section is Algorithm 3 used to compute such strategies as well as the vector $\text{SafePR}_{\mathcal{M}, T}$ of minimal initial resource levels for which such a strategy exist. As before, for the rest of this section we fix a CMDP \mathcal{M} .

We define a function $\text{SPR-Val}_{\mathcal{M}}: S \times A \times \overline{\mathbb{N}}^S \rightarrow \overline{\mathbb{N}}$ (*SPR* for safe positive reachability) s.t. for all $s \in S, a \in A$, and $\mathbf{x} \in \overline{\mathbb{N}}^S$ we have

$$\text{SPR-Val}_{\mathcal{M}}(s, a, \mathbf{x}) = C(s, a) + \min_{t \in \text{Succ}(s, a)} \left\{ \max \{ \mathbf{x}(t), \text{Safe}_{\mathcal{M}}(t') \mid t' \in \text{Succ}(s, a), t' \neq t \} \right\}$$

The max operator considers, for given t , the value $\mathbf{x}(t)$ and the values needed to survive from all possible outcomes of a other than t . Let $v = \text{SPR-Val}_{\mathcal{M}}(s, a, \mathbf{x})$ and t the outcome selected by min. Intuitively, v is the minimal amount of resource needed to reach t with at least $\mathbf{x}(t)$ resource units, or survive if the outcome of a is different from t .

We now define a functional whose fixed point characterizes $\text{SPR-Val}_{\mathcal{M}, T}$. We first define a two-sided version of the truncation operator from the previous section: the operator $\llbracket \cdot \rrbracket_{\mathcal{M}}$ such that

$$\llbracket \mathbf{x} \rrbracket_{\mathcal{M}}(s) = \begin{cases} \infty & \text{if } \mathbf{x}(s) > \text{cap}(\mathcal{M}) \\ \mathbf{x}(s) & \text{if } \mathbf{x}(s) \leq \text{cap}(\mathcal{M}) \text{ and } s \notin R \\ 0 & \text{if } \mathbf{x}(s) \leq \text{cap}(\mathcal{M}) \text{ and } s \in R \end{cases}$$

Using the functions $SPR\text{-}Val$ and $\llbracket \cdot \rrbracket_M$, we now define an auxiliary operator \mathcal{A} and the main operator \mathcal{B} as follows.

$$\mathcal{A}_M(\mathbf{r})(s) = \begin{cases} Safe_M(s) & \text{if } s \in T \\ \min_{a \in A} (SPR\text{-}Val_M(s, a, \mathbf{r})) & \text{otherwise;} \end{cases}$$

$$\mathcal{B}_M(\mathbf{r}) = \llbracket \mathcal{A}_M(\mathbf{r}) \rrbracket_M$$

Let $SafePR_T^i$ be the vector such that for a state $s \in S$ the number $d = SafePR_T^i(s)$ is the minimal number ℓ such that there exists a strategy that is ℓ -safe in s and produces at least one run that visits T within first i steps. Further, we denote by \mathbf{y}_T a vector such that

$$\mathbf{y}_T(s) = \begin{cases} Safe_M(s) & \text{if } s \in T \\ \infty & \text{if } s \notin T \end{cases}$$

The following lemma is proved by a rather straightforward but technical induction. The proof is deferred to Appendix A.4.

Lemma 4. *Consider the iteration of \mathcal{B}_M on the initial vector \mathbf{y}_T . Then for each $i \geq 0$ it holds that $\mathcal{B}_M^i(\mathbf{y}_T) = SafePR_{M,T}^i$.*

The following lemma says that iterating \mathcal{B}_M reaches a fixed point in a polynomial number of iterations. Intuitively, this is because when trying to reach T , it doesn't make sense to perform a cycle between two visits of a reload state (as this can only increase the resource consumption) and at the same time it doesn't make sense to visit the same reload state twice (since the resource is reloaded to the full capacity upon each visit). The proof is deferred to Appendix A.5.

Lemma 5. *Let $K = |R| + (|R| + 1) \cdot (|S| - |R| + 1)$. Taking the same initial vector \mathbf{y}_T as in Lemma 4, we have $\mathcal{B}_M^K(\mathbf{y}_T) = SafePR_{M,T}$.*

The computation of $SafePR_{M,T}$ and of the associated witness strategy is presented in Algorithm 3.

Theorem 6. *The Algorithm 3 always terminates after a polynomial number of steps, and upon termination, $\mathbf{r} = SafePR_{M,T}$.*

Proof. The repeat loop on lines 1–4 initialize \mathbf{r} to \mathbf{y}_T . The repeat loop on lines 5–14 then iterates the operator \mathcal{B} . By Lemma 5, the iteration reaches a fixed point in at most K steps, and this fixed point equals $SafePR_{M,T}$. The complexity bound follows easily, since K is of polynomial magnitude.

The most intricate part of our analysis is extracting a strategy that is T -positive $SafePR_{M,T}(s)$ -safe in every state s .

Theorem 7. *Let $\mathbf{v} = SafePR_{M,T}$. Upon termination of Algorithm 3, the computed selector Σ has the property that the finite counter strategy $\Sigma^{\mathbf{v}}$ is, for each state $s \in S$, T -positive $\mathbf{v}(s)$ -safe in s . That is, a polynomial-size finite counter strategy for the positive reachability problem can be computed in polynomial time.*

Algorithm 3: Positive reachability of T in \mathcal{M}

Input: CMDP \mathcal{M} with states S , set of target states $T \subseteq S$
Output: The vector $\text{SafePR}_{\mathcal{M},T}$, coreresponding rule selector Σ

```
1  $\mathbf{r} \leftarrow \{\infty\}^S$ ;  
2 foreach  $s \in S$  s.t.  $\text{Safe}_{\mathcal{M}}(s) < \infty$  do  
3    $\Sigma(s)(\text{Safe}_{\mathcal{M}}(s)) \leftarrow$  arbitrary action safe in  $s$   
4 foreach  $t \in T$  do  $\mathbf{r}(t) \leftarrow \text{Safe}_{\mathcal{M}}(t)$  ;  
5 repeat  
6    $\mathbf{r}_{old} \leftarrow \mathbf{r}$ ;  
7   foreach  $s \in S \setminus T$  do  
8      $\mathbf{a}(s) \leftarrow \arg \min_{a \in A} \text{SPR-Val}(s, a, \mathbf{r}_{old})$ ;  
9      $\mathbf{r}(s) \leftarrow \min_{a \in A} \text{SPR-Val}(s, a, \mathbf{r}_{old})$ ;  
10   $\mathbf{r} \leftarrow \llbracket \mathbf{r} \rrbracket_{\mathcal{M}}$ ;  
11  foreach  $s \in S \setminus T$  do  
12    if  $\mathbf{r}(s) < \mathbf{r}_{old}(s)$  then  
13       $\Sigma(s)(\mathbf{r}(s)) \leftarrow \mathbf{a}(s)$ ;  
14 until  $\mathbf{r}_{old} = \mathbf{r}$ ;  
15 return  $\mathbf{r}, \Sigma$ 
```

The rest of this section is devoted to the proof of Theorem 7. The complexity follows from Theorem 6. Indeed, since the algorithm has a polynomial complexity, also the size of Σ is polynomial.

The correctness proof is based on the following invariant of the main repeat loop: the finite counter strategy $\pi = \Sigma^{\mathbf{r}}$ has these properties:

- (a) We have that π is $\text{Safe}_{\mathcal{M}}(s)$ -safe in every state $s \in S$; in particular, we have for $l = \min\{\mathbf{r}(s), \text{cap}(\mathcal{M})\}$ that $RL_l(\alpha) \neq \perp$ for every finite path α produced by π from s .
- (b) For each state $s \in S$ such that $\mathbf{r}(s) \leq \text{cap}(\mathcal{M})$ there exists a π -compatible finite path $\alpha = s_1 a_1 s_2 \dots s_n$ such that $s_1 = s$ and $s_n \in T$ and such that “the resource level with initial load $\mathbf{r}(s)$ never decreases below \mathbf{r} along α ”, which means that for each prefix $\alpha_{..i}$ of α it holds $RL_{\mathbf{r}(s)}(\alpha_{..i}) \geq \mathbf{r}(s_i)$.

The theorem then follows from this invariant (parts (a) and the first half of (b)) and from Theorem 6. We start with the following support invariant, which is easy to prove.

Lemma 6. *The inequality $\mathbf{r} \geq \text{Safe}_{\mathcal{M}}$ is an invariant of the main repeat-loop.*

Proving part (a) of the main invariant. We use the following auxiliary lemma.

Lemma 7. *Assume that Σ is a counter selector such that for all $s \in S$ such that $\text{Safe}(s) < \infty$:*

- (1.) $\text{Safe}(s) \in \text{dom}(\Sigma(s))$.
- (2.) *For all $x \in \text{dom}(\Sigma(s))$, for $a = \Sigma(s)(x)$ and for all $t \in \text{Succ}(s, a)$ we have $RL_x(\text{sat}) = d - C(s, a) \geq \text{Safe}(t)$ where $d = x$ for $s \notin R$ and $d = \text{cap}(\mathcal{M})$ otherwise.*

Then for each vector $\mathbf{y} \geq \text{Safe}$ the strategy $\pi = \Sigma^{\mathbf{y}}$ is $\text{Safe}(s)$ -safe in every state s .

Proof. Let s be a state such that $\mathbf{y}(s) < \infty$. It suffices to prove that for every π -compatible finite path α started in s it holds $\perp \neq RL_{\mathbf{y}(s)}(\alpha)$. We actually prove a stronger statement: $\perp \neq RL_{\mathbf{y}(s)}(\alpha) \geq \text{Safe}(\text{last}(\alpha))$. We proceed by induction on the length of α . If $\text{len}(\alpha) = 0$ we have $RL_{\mathbf{y}(s)}(\alpha) = \mathbf{y}(s) \geq \text{Safe}_{\mathcal{M}}(s) \geq 0$. Now let $\alpha = \beta \odot t_1 at_2$ for some shorter path β with $\text{last}(\beta) = t_1$ and $a \in A$, $t_1, t_2 \in S$. By induction hypothesis, $l = RL_{\mathbf{y}(s)}(\beta) \geq \text{Safe}_{\mathcal{M}}(t_1)$, from which it follows that $\text{Safe}_{\mathcal{M}}(t_1) < \infty$. Due to (1.), it follows that there exists at least one $x \in \text{dom}(\Sigma(t_1))$ such that $x \leq l$. We select maximal x satisfying the inequality so that $a = \Sigma(t_1)(x)$. We have that $RL_{\mathbf{y}(s)}(\alpha) = RL_l(t_1 at_2)$ by definition and from (2.) it follows that $\perp \neq RL_x(t_1 at_2) \geq \text{Safe}(t_2) \geq 0$. All together, as $l \geq x$ we have that $RL_{\mathbf{y}(s)}(\alpha) \geq RL_x(t_1 at_2) \geq \text{Safe}(t_2) \geq 0$. \square

Now we prove the part (a) of the main invariant. We show that throughout the execution of Algorithm 3, Σ satisfies the assumptions of Lemma 7. Property (1.) is ensured by the initialization on line 3. The property (2.) holds upon first entry to the main loop by the definition of a safe action (Definition 6). Now assume that $\Sigma(s)(\mathbf{r}(s))$ is redefined on line 13, and let a be the action $\mathbf{a}(s)$.

We first handle the case when $s \notin R$. Since a was selected on line 8, from the definition of *SPR-Val* we have that there is $t \in \text{Succ}(s, a)$ such that after the loop iteration,

$$\mathbf{r}(s) = C(s, a) + \max\{\mathbf{r}_{old}(t), \text{Safe}(t') \mid t \neq t' \in \text{Succ}(s, a)\} \geq C(s, a) + \max_{t' \in \text{Succ}(s, a)} \text{Safe}_{\mathcal{M}}(t'), \quad (2)$$

the latter inequality following from Lemma 6. Satisfaction of property (2.) in s then follows immediately from the equation (2).

If $s \in R$, then (2) holds before the truncation on line 10, at which point $\mathbf{r}(s) < \text{cap}(\mathcal{M})$. Hence, $\text{cap}(\mathcal{M}) - C(s, a) \geq \max_{t \in \text{Succ}(s, a)} \text{Safe}_{\mathcal{M}}(t)$ as required by (2.). From Lemmas 6 and 7 it follows that $\Sigma^{\mathbf{r}}$ is $\text{Safe}_{\mathcal{M}}(s)$ -safe in every state s . This finishes the proof of part (a) of the invariant.

Proving part (b) of the main invariant. Clearly, (b) holds right after initialization. Now assume that an iteration of the main repeat loop was performed. Denote π_{old} denote the strategy $\Sigma^{\mathbf{r}_{old}}$ and by π the strategy $\Sigma^{\mathbf{r}}$. Let s be any state such that $\mathbf{r}(s) \leq \text{cap}(\mathcal{M})$. If $\mathbf{r}(s) = \mathbf{r}_{old}(s)$, then we claim that (b) follows directly from the induction hypothesis: indeed, by induction hypothesis we have that there is an s -initiated π_{old} -compatible path α ending in a target state s.t. the $\mathbf{r}_{old}(s)$ -initiated resource level along α never drops \mathbf{r}_{old} , i.e. for each prefix β of α it holds $RL_{\mathbf{r}_{old}(s)}(\beta) \geq \mathbf{r}_{old}(\text{last}(\beta))$. But then β is also π -compatible, since for each state q , $\Sigma(q)$ was only redefined for values smaller than $\mathbf{r}_{old}(q)$.

The case when $\mathbf{r}(s) < \mathbf{r}_{old}(s)$ is treated similarly. As in the proof of part (a), denote by a the action $\mathbf{a}(s)$ assigned on line 13. There must be a state $t \in \text{Succ}(s, a)$ s.t. (2) holds before the truncation on line 10. In particular, for this t it holds $RL_{\mathbf{r}(s)}(\text{sat}) \geq \mathbf{r}_{old}(t)$. By induction hypothesis, there is a t -initiated π_{old} -compatible path β ending in T satisfying the conditions in (b). We put $\alpha = \text{sat} \odot \beta$. Clearly α is s -initiated and reaches T . Moreover, it is π -compatible. To see this, note that $\Sigma^{\mathbf{r}}(s)(\mathbf{r}(s)) = a$; moreover, the resource level after the first transition is $e(t) = RL_{\mathbf{r}(s)}(\text{sat}) \geq \mathbf{r}_{old}(t)$, and due to the

assumed properties of β , the $\mathbf{r}_{old}(t)$ -initiated resource level (with initial load $e(t)$) never decreases below \mathbf{r}_{old} along β . Since Σ was only re-defined for values smaller than those given by the vector \mathbf{r}_{old} , π mimics π_{old} along β . Since $\mathbf{r} \leq \mathbf{r}_{old}$, we have that along α , the $\mathbf{r}(s)$ -initiated resource level never decreases below \mathbf{r} . This finishes the proof of part (b) of the invariant and thus also the proof of Theorem 7 \square

6 Büchi

This section proves Theorem 1 which is the main theoretical result of the paper. The proof is broken down into the following steps.

- (1.) We identify a largest set $R' \subseteq R$ of reload states such that from each $r \in R'$ we can reach R' again (in at least one step) while consuming at most cap resource units and restricting ourselves only to strategies that (i) avoid $R \setminus R'$ and (ii) guarantee positive reachability of T in $\mathcal{M}(R')$.
- (2.) We show that $SafeBüchi_{\mathcal{M},T} = SafePR_{\mathcal{M}(R'),T}$ and that the corresponding strategy (computed by Algorithm 3) is also T -Büchi $SafeBüchi_{\mathcal{M},T}(s)$ -safe for each $s \in S$.

Algorithm 4 solves (1.) in a similar fashion as Algorithm 2 handled safety. In each iteration, we declare as non-reloading all states from which positive reachability of T and safety within $\mathcal{M}(Rel)$ cannot be guaranteed. This is repeated until we reach a fixed point. The number of iterations is clearly bounded by $|R|$.

Algorithm 4: Almost-sure Büchi reachability of T in \mathcal{M} .

Input: CMDP $\mathcal{M} = (S, A, \Delta, C, R, cap)$, target states $T \subseteq S$
Output: The largest set $Rel \subseteq R$ such that $SafePR_{\mathcal{M}(Rel),T}(r) \leq cap$ for all $r \in Rel$.

```

1  $Rel \leftarrow R; ToRemove \leftarrow \emptyset;$ 
2 repeat
3    $Rel \leftarrow Rel \setminus ToRemove;$ 
4    $(\mathbf{reach}, \Sigma) \leftarrow SafePR_{\mathcal{M}(Rel),T};$ 
5    $ToRemove \leftarrow \{r \in Rel \mid \mathbf{reach}(r) > cap\};$ 
6 until  $ToRemove = \emptyset;$ 
7 return  $\mathbf{reach}, \Sigma$ 
```

Theorem 8. Let $\mathcal{M} = (S, A, \Delta, C, R, cap)$ be a CMDP and $T \subseteq S$ be a target set. Moreover, let R' be the contents of Rel upon termination of Algorithm 4 for the input \mathcal{M} and T . Finally let \mathbf{r} and Σ be the vector and the selector returned by Algorithm 3 for the input \mathcal{M} and T . Then for every state s , the finite counter strategy $\sigma = \Sigma^{\mathbf{r}}$ is T -Büchi $\mathbf{r}(s)$ -safe in s in both $\mathcal{M}(R')$ and \mathcal{M} . Moreover, the vector \mathbf{r} is equal to $SafeBüchi_{\mathcal{M},T}$.

Proof. We first show that σ is T -Büchi $\mathbf{r}(s)$ -safe in $\mathcal{M}(R')$ for all $s \in S$ with $\mathbf{r}(s) \leq cap$. Clearly it is $\mathbf{r}(s)$ -safe, so it remains to prove that T is visited infinitely often with

probability 1. We know that upon every visit of a state $r \in R'$, σ guarantees a future visit to T with positive probability. As a matter of fact, since σ is a finite memory strategy, there is $\delta > 0$ such that upon every visit of some $r \in R'$, the probability of a future visit to T is at least δ . As $M(R')$ is decreasing, every s -initiated σ -compatible run must visit the set R' infinitely many times. Hence, with probability 1 we reach T at least once. The argument can then be repeated from the first point of visit to T to show that with probability 1 we visit T at least twice, three times, etc. *ad infinitum*. By the monotonicity of probability, $\mathbb{P}_{M,s}^\sigma(\text{Büchi}_T) = 1$.

It remains to show that $\mathbf{r} \leq \text{SafeBüchi}_{M,T}$. Assume that there is a state $s \in S$ and a strategy σ' such that σ' is d -safe in s for some $d < \mathbf{r}(s) = \text{SafePR}_{M(R'),T}(s)$. We show that this strategy is not T -Büchi d -safe in M . If all σ' -compatible runs reach T , then there must be at least one history α produced by σ' that visits $r \in R \setminus R'$ before reaching T (otherwise $d \geq \mathbf{r}(s)$). Then either (a) $\text{SafePR}_{M,T}(r) = \infty$, in which case any σ' -compatible extension of α avoids T ; or (b) since $\text{SafePR}_{M(R'),T}(r) > \text{cap}$, there must be an extension of α that visits, between the visit of r and T , another $r' \in R \setminus R'$ such that $r' \neq r$. We can then repeat the argument, eventually reaching the case (a) or running out of the resource, a contradiction with σ' being d -safe. \square

We can finally proceed to prove Theorem 1.

Proof (of Theorem 1). The theorem follows immediately from Theorem 8 since we can (1.) compute $\text{SafeBüchi}_{M,T}$ and the corresponding strategy σ_T in polynomial time (see Theorem 7 and Algorithm 4), (2.) we can easily check whether $d \geq \text{SafeBüchi}_{M,T}(s)$, if yes, then σ_T is the desired strategy σ and (3.) represent σ_T in polynomial space as it is a finite counter strategy represented by a polynomial-size counter selector. \square

7 Implementation and Case Studies

We have implemented the presented algorithms in Python in a tool called *FiMDP* (*Fuel in MDP*) available at <https://github.com/xblahoud/FiMDP>. The docker artifact is available at <https://hub.docker.com/r/xblahoud/fimdp> and can be run without installation via the Binder project [53]. Integration of the tool with the interactive Jupyter notebooks [48] allows users to investigate the behavior of the algorithms visually on small examples visually. The artifact contains notebooks discussing interesting examples and incorrect approaches to solving the problems.

We investigate the practical behavior of our algorithms using two case studies: (1) An autonomous electric vehicle (AEV) routing problem modeled using realistic traffic and electric car energy consumption data, and (2) a multi-agent grid world model inspired by the Mars Helicopter Scout [8] to be deployed from the planned Mars 2020 rover. We use the first scenario to show the utility of our algorithm for solving real-world problems and the second scenario to benchmark the algorithm's scalability.

7.1 Electric Vehicle Routing

To demonstrate the utility of our approach, the first case study discusses the routing of an AEV on the streets of Manhattan. Due to their limited driving range and availability of recharge stations, AEV routing presents a significant challenge in planning [60].

We consider an area in the middle of Manhattan, from 42nd to 116th Street, given in Fig. 2, with the agent’s original state space consisting of street intersections. The agent’s actions — turns performed at these intersections — result in deterministic transitions to subsequent intersections, but with stochastic consumption of energy, depending on traffic. We model the consumption on each edge by using data on distributions of vehicle travel times from the area [51] and conversion of velocity and travel times to energy consumption [55]. We then discretize the energy consumption distribution into D possible consumptions c_1, \dots, c_D (with sets $\{c_1, \dots, c_D\}$ possibly different at different edges). Since the CMDP framework allows only for deterministic consumption, we encode the stochastic consumption in the following way: instead of connecting each two neighboring intersections s and s' directly, they are connected by edges from s to d_i and d_i to s' , where d_1, \dots, d_D are *dummy states*. The probability of transitioning from s to d_i with consumption c_i is set to the probability that the consumption equals c_i . Playing any action in d_i leads to s' with probability 1 and consumption 0. Intersections in the proximity of real-world fast charging stations [52] represent the set of reload states.




Fig. 2. Visualization of the street network in the area considered for our study. Green edges represent one-way roads, blue edges represent two-way roads, and red nodes represent reload states.

Modifying the graph from Fig. 2 as described above with $D = 3$, we obtain a CMDP with 7378 states. We randomly select 100 nodes as destinations T and run the presented algorithms on a Lenovo IdeaPad 330S laptop with 8GB RAM. The mean times (7 iterations, standard deviations less than 5% in every case) needed to compute the vectors Safe , SafePR_T , and SafeBüchi_T in seconds are: 2.21, 3.02, and 3.11, respectively. We observe that, despite the large number of states, our implementation of the proposed algorithms runs in a reasonable time. Moreover, Table 1 shows that restricting the capacity decreases the running times even further.

Table 1. Mean computation times of SafeBüchi_T depending on AEV energy capacity

Capacity	3	5	7	10	15	20	100
mean running time [sec]	0.29	0.57	0.93	1.32	1.51	2.36	3.11

7.2 Multi-agent Grid World

Our second case study aims to reach the scalability limits of the proposed algorithms. It explores a multi-agent scenario of a rover and a helicopter operating on Mars. The narrative of the case study is informed by realistic considerations of the Mars 2020 mission. Namely, we consider a rover of infinite energy capacity and a helicopter of finite capacity that recharges by returning to the rover. These two vehicle jointly operate on a mission where the helicopter needs to reach areas inaccessible to the rover. We assume that the outcomes of all the actions of the helicopter are deterministic while those of the rover — influenced by terrain dynamics — are stochastic. For the purpose of this experiment, we specify that every transition of the helicopter costs 1 energy unit. For a grid world of size n , this system can be naturally modeled as a CMDP with n^4 states (representing the x - and y -coordinates of the rover and the helicopter).

As in the previous case study, we first calculate the computation time for the three different objectives. On a grid world of size $n = 10$ (resulting in a MDP with 10000 states), with the helicopter energy capacity of 10 units, we obtain the following mean computation times: safety – 17.3 sec, positive reachability – 19.3 sec, and Büchi – 19.3 sec. Table 2 shows mean computational time (with standard deviation not exceeding 2.5%) for the the Büchi objective and growing grid sizes.

Table 2. Computation time of *SafeBüchi_T* depending on grid size

Grid size	5	7	10	15	20
Number of states	625	2401	10000	50625	160000
Mean (sec)	0.619	3.85	19.3	106	363

We observe that the increase in the mean computation time follows the growth in the number of states roughly linearly, and our implementation deals with an MDP with 1.6×10^5 states in only six minutes.

8 Conclusion & Future Work

We presented a first study of consumption Markov decision processes (CMDPs) with qualitative ω -regular objectives. We developed and implemented a polynomial-time algorithm for CMDPs with an objective of probability-1 satisfaction of a given Büchi condition. Possible directions for the future work are extensions to quantitative analysis (e.g. minimizing the expected resource consumption), stochastic games, or partially observable setting.

Acknowledgement

Petr Novotný is supported by the Czech Science Foundation grant no. GJ19-15134Y. This work was partly supported by an Early Stage Innovations grant from NASA’s Space Technology Research Grants Program, grant no. 80NSSC19K0209 and by DARPA grant no. HR001120C0065.

References

1. P. A. Abdulla, M. F. Atig, P. Hofman, R. Mayr, K. N. Kumar, and P. Totzke. Infinite-state energy games. In *Joint Meeting of the Twenty-Third EACSL Annual Conference on Computer Science Logic (CSL) and the Twenty-Ninth Annual ACM/IEEE Symposium on Logic in Computer Science (LICS), CSL-LICS '14, Vienna, Austria, July 14 - 18, 2014*, pages 7:1–7:10, 2014.
2. R. Ash and C. Doléans-Dade. *Probability and Measure Theory*. Harcourt/Academic Press, 2000.
3. G. Bacci, P. Bouyer, U. Fahrenberg, K. G. Larsen, N. Markey, and P.-A. Reynier. Optimal and robust controller synthesis. In K. Havelund, J. Peleska, B. Roscoe, and E. de Vink, editors, *Formal Methods*, pages 203–221, Cham, 2018. Springer International Publishing.
4. C. Baier, P. Chrszon, C. Dubslaff, J. Klein, and S. Klüppelholz. Energy-utility analysis of probabilistic systems with exogenous coordination. In F. de Boer, M. Bonsangue, and J. Rutten, editors, *It's All About Coordination: Essays to Celebrate the Lifelong Scientific Achievements of Farhad Arbab*, pages 38–56. Springer International Publishing, Cham, 2018.
5. C. Baier, M. Daum, C. Dubslaff, J. Klein, and S. Klüppelholz. Energy-utility quantiles. In J. M. Badger and K. Y. Rozier, editors, *NASA Formal Methods*, pages 285–299, Cham, 2014. Springer International Publishing.
6. C. Baier, C. Dubslaff, J. Klein, S. Klüppelholz, and S. Wunderlich. Probabilistic model checking for energy-utility analysis. In F. van Breugel, E. Kashefi, C. Palamidessi, and J. Rutten, editors, *Horizons of the Mind. A Tribute to Prakash Panangaden: Essays Dedicated to Prakash Panangaden on the Occasion of His 60th Birthday*, pages 96–123. Springer International Publishing, Cham, 2014.
7. C. Baier, C. Dubslaff, S. Klüppelholz, and L. Leuschner. Energy-utility analysis for resilient systems using probabilistic model checking. In G. Ciardo and E. Kindler, editors, *Application and Theory of Petri Nets and Concurrency*, pages 20–39, Cham, 2014. Springer International Publishing.
8. B. Balaram, T. Canham, C. Duncan, H. F. Grip, W. Johnson, J. Maki, A. Quon, R. Stern, and D. Zhu. Mars helicopter technology demonstrator. In *2018 AIAA Atmospheric Flight Mechanics Conference*, page 0023, 2018.
9. U. Boker, T. A. Henzinger, and A. Radhakrishna. Battery transition systems. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 595–606. ACM, 2014.
10. P. Bouyer, U. Fahrenberg, K. G. Larsen, and N. Markey. Timed automata with observers under energy constraints. In *Proceedings of the 13th ACM international conference on Hybrid systems: computation and control*, pages 61–70. ACM, 2010.
11. P. Bouyer, U. Fahrenberg, K. G. Larsen, N. Markey, and J. Srba. Infinite Runs in Weighted Timed Automata with Energy Constraints. In F. Cassez and C. Jard, editors, *Proceedings of FORMATS 2008*, volume 5215 of *LNCS*, pages 33–47. Springer Berlin Heidelberg, 2008.
12. P. Bouyer, P. Hofman, N. Markey, M. Randour, and M. Zimmermann. Bounding average-energy games. In J. Esparza and A. S. Murawski, editors, *Foundations of Software Science and Computation Structures*, pages 179–195, Berlin, Heidelberg, 2017. Springer Berlin Heidelberg.
13. P. Bouyer, N. Markey, M. Randour, K. G. Larsen, and S. Laursen. Average-energy games. *Acta Informatica*, 55(2):91–127, Mar 2018.
14. T. Brázdil, K. Chatterjee, A. Kučera, and P. Novotný. Efficient Controller Synthesis for Consumption Games with Multiple Resource Types. In *Proceedings of CAV 2012*, volume 7358 of *LNCS*, pages 23–38. Springer, 2012.

15. T. Brázdil, P. Jančar, and A. Kučera. Reachability Games on Extended Vector Addition Systems with States. In S. Abramsky, C. Gavaille, C. Kirchner, and F. Meyer auf der Heide, editors, *Automata, Languages and Programming*, volume 6199 of *LNCS*, pages 478–489, 2010.
16. T. Brázdil, D. Klaška, A. Kučera, and P. Novotný. Minimizing Running Costs in Consumption Systems. In A. Biere and R. Bloem, editors, *Computer Aided Verification 2014*, volume 8559 of *LNCS*, pages 457–472. Springer International Publishing, 2014.
17. T. Brázdil, A. Kucera, and P. Novotný. Optimizing the expected mean payoff in energy Markov decision processes. In *Automated Technology for Verification and Analysis - 14th International Symposium, ATVA 2016, Chiba, Japan, October 17-20, 2016, Proceedings*, pages 32–49, 2016.
18. R. Brenguier, F. Cassez, and J.-F. Raskin. Energy and mean-payoff timed games. In *Proceedings of the 17th International Conference on Hybrid Systems: Computation and Control, HSCC '14*, page 283–292, New York, NY, USA, 2014. Association for Computing Machinery.
19. T. Brihaye, G. Geeraerts, A. Haddad, and B. Monmege. Pseudopolynomial iterative algorithm to solve total-payoff games and min-cost reachability games. *Acta Informatica*, 54(1):85–125, Feb 2017.
20. L. Brim, J. Chaloupka, L. Doyen, R. Gentilini, and J. Raskin. Faster algorithms for mean-payoff games. *Formal Methods in System Design*, 38(2):97–118, 2011.
21. V. Bruyère, Q. Hautem, M. Randour, and J.-F. Raskin. Energy Mean-Payoff Games. In W. Fokink and R. van Glabbeek, editors, *30th International Conference on Concurrency Theory (CONCUR 2019)*, volume 140 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 21:1–21:17, Dagstuhl, Germany, 2019. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik.
22. D. Cachera, U. Fahrenberg, and A. Legay. An ω -Algebra for Real-Time Energy Problems. *Logical Methods in Computer Science*, Volume 15, Issue 2, May 2019.
23. A. Chakrabarti, L. de Alfaro, T. A. Henzinger, and M. Stoelinga. Resource Interfaces. In R. Alur and I. Lee, editors, *Proceedings of EMSOFT 2003*, volume 2855 of *LNCS*, pages 117–133, Heidelberg, 2003. Springer.
24. J. Chaloupka. Z-reachability problem for games on 2-dimensional vector addition systems with states is in P. *Fundam. Inform.*, 123(1):15–42, 2013.
25. K. Chatterjee. *Stochastic ω -regular games*. PhD thesis, University of California, Berkeley, 2007.
26. K. Chatterjee and L. Doyen. Energy and Mean-Payoff Parity Markov Decision Processes. In *Proceedings of MFCS 2011*, volume 6907 of *LNCS*, pages 206–218. Springer, 2011.
27. K. Chatterjee and L. Doyen. Energy parity games. *Theoretical Computer Science*, 458:49–60, 2012.
28. K. Chatterjee, L. Doyen, T. Henzinger, and J.-F. Raskin. Generalized Mean-payoff and Energy Games. In K. Lodaya and M. Mahajan, editors, *Proceedings of IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science (FST&TCS 2010)*, volume 8, pages 505–516, 2010.
29. K. Chatterjee, M. Henzinger, S. Krinninger, and D. Nanongkai. Polynomial-Time Algorithms for Energy Games with Special Weight Structures. In *Proceedings of ESA 2012*, pages 301–312. Springer:, series = LNCS, title =.
30. K. Chatterjee, M. Jurdziński, and T. Henzinger. Quantitative stochastic parity games. In *Proceedings of SODA 2004*, pages 121–130, 2004.
31. K. Chatterjee, M. Randour, and J.-F. Raskin. Strategy synthesis for multi-dimensional quantitative objectives. *Acta informatica*, 51(3-4):129–163, 2014.
32. C. Courcoubetis and M. Yannakakis. The Complexity of Probabilistic Verification. *J. ACM*, 42(4):857–907, 1995.

33. L. de Alfaro. *Formal verification of probabilistic systems*. Phd. thesis, Stanford University, Stanford, CA, USA, 1998.
34. A. Degorre, L. Doyen, R. Gentilini, J.-F. Raskin, and S. Toruńczyk. Energy and mean-payoff games with imperfect information. In *Computer Science Logic*, pages 260–274, 2010.
35. Z. Ésik, U. Fahrenberg, A. Legay, and K. Quaas. An algebraic approach to energy problems I - continuous Kleene ω -algebras. *Acta Cybernetica*, 23(1):203–228, Jan. 2017.
36. Z. Ésik, U. Fahrenberg, A. Legay, and K. Quaas. An algebraic approach to energy problems II - the algebra of energy functions. *Acta Cybernetica*, 23(1):229–268, Jan. 2017.
37. U. Fahrenberg, L. Juhl, K. G. Larsen, and J. Srba. Energy Games in Multiweighted Automata. In A. Cerone and P. Pihlajasaari, editors, *Theoretical Aspects of Computing – ICTAC 2011*, volume 6916 of *LNCS*, pages 95–115, Heidelberg, 2011. Springer.
38. U. Fahrenberg and A. Legay. Featured weighted automata. In *2017 IEEE/ACM 5th International FME Workshop on Formal Methods in Software Engineering (FormaliSE)*, pages 51–57, May 2017.
39. N. Fijalkow and M. Zimmermann. Cost-Parity and Cost-Streett Games. In D. D’Souza, T. Kavitha, and J. Radhakrishnan, editors, *IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science (FSTTCS 2012)*, volume 18 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 124–135, Dagstuhl, Germany, 2012. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik.
40. E. Filiot, R. Gentilini, and J.-F. Raskin. Quantitative languages defined by functional automata. In M. Koutny and I. Ulidowski, editors, *CONCUR 2012 – Concurrency Theory*, pages 132–146, Berlin, Heidelberg, 2012. Springer Berlin Heidelberg.
41. E. M. Hahn, M. Perez, F. Somenzi, A. Trivedi, S. Schewe, and D. Wojtczak. Good-for-MDPs automata for probabilistic analysis and reinforcement learning. *CoRR*, abs/1909.05081. To appear in TACAS’20., 2019.
42. L. Herrmann, C. Baier, C. Fetzer, S. Klüppelholz, and M. Napierkowski. Formal parameter synthesis for energy-utility-optimal fault tolerance. In R. Bakhshi, P. Ballarini, B. Barbot, H. Castel-Taleb, and A. Remke, editors, *Computer Performance Engineering*, pages 78–93, Cham, 2018. Springer International Publishing.
43. L. Juhl, K. G. Larsen, and J.-F. Raskin. Optimal bounds for multiweighted and parametrised energy games. In *Theories of Programming and Formal Methods*, pages 244–255. Springer, 2013.
44. M. Jurdziński, R. Lazić, and S. Schmitz. Fixed-dimensional energy games are in pseudo-polynomial time. In M. M. Halldórsson, K. Iwama, N. Kobayashi, and B. Speckmann, editors, *Automata, Languages, and Programming*, pages 260–272, Berlin, Heidelberg, 2015. Springer Berlin Heidelberg.
45. M. Jurdziński. Deciding the winner in parity games is in $UP \cap co-UP$. *Information Processing Letters*, 68(3):119 – 124, 1998.
46. L. Khachiyan, E. Boros, K. Borys, K. Elbassioni, V. Gurvich, G. Rudolf, and J. Zhao. On short paths interdiction problems: Total and node-wise limited interdiction. *Theory of Computing Systems*, 43(2):204–233, 2008.
47. D. Klačka. *Complexity of Consumption Games*. Bachelor’s thesis, Masaryk University, 2014.
48. T. Kluyver, B. Ragan-Kelley, F. Pérez, B. Granger, M. Bussonnier, J. Frederic, K. Kelley, J. Hamrick, J. Grout, S. Corlay, P. Ivanov, D. Avila, S. Abdalla, C. Willing, and Jupyter development team. Jupyter notebooks — a publishing format for reproducible computational workflows. In F. Loizides and B. Schmidt, editors, *Proceedings of the 20th International Conference on Electronic Publishing: Positioning and Power in Academic Publishing: Players, Agents and Agendas (ELPUB’16)*, pages 87–90. IOS Press, 2016.
49. K. G. Larsen, S. Laursen, and M. Zimmermann. Limit your consumption! Finding bounds in average-energy games. In *Proceedings of the 14th International Workshop Quantitative*

- Aspects of Programming Languages and Systems, QAPL 2016, Eindhoven, The Netherlands, April 2-3, 2016*, pages 1–14, 2016.
50. R. Mayr, S. Schewe, P. Totzke, and D. Wojtczak. MDPs with energy-parity objectives. In *2017 32nd Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*, pages 1–12, June 2017.
 51. U. Movement. *Traffic Speed Data for New York City*, 2019.
 52. U. S. D. of Energy. Alternative fuels data center, 2019.
 53. Project Jupyter, Matthias Bussonnier, Jessica Forde, Jeremy Freeman, Brian Granger, Tim Head, Chris Holdgraf, Kyle Kelley, Gladys Nalvarte, Andrew Osherooff, M. Pacer, Yuvi Panda, Fernando Perez, Benjamin Ragan Kelley, and Carol Willing. Binder 2.0 - Reproducible, interactive, sharable environments for science at scale. In Fatih Akici, David Lippa, Dillon Niederhut, and M. Pacer, editors, *Proceedings of the 17th Python in Science Conference*, pages 113 – 120, 2018.
 54. S. Sickert, J. Esparza, S. Jaax, and J. Křetínský. Limit-deterministic büchi automata for linear temporal logic. In S. Chaudhuri and A. Farzan, editors, *Computer Aided Verification*, pages 312–332, Cham, 2016. Springer International Publishing.
 55. J. Straubel. *Roadster Efficiency and Range*, 2008.
 56. G. Sugumar, R. Selvamuthukumar, T. Dragicevic, U. Nyman, K. G. Larsen, and F. Blaabjerg. Formal validation of supervisory energy management systems for microgrids. In *IECON 2017 - 43rd Annual Conference of the IEEE Industrial Electronics Society*, pages 1154–1159, Oct 2017.
 57. R. S. Sutton and A. G. Barto. *Reinforcement learning: An introduction*. MIT press, 2018.
 58. Y. Velner, K. Chatterjee, L. Doyen, T. A. Henzinger, A. M. Rabinovich, and J. Raskin. The complexity of multi-mean-payoff and multi-energy games. *Inf. Comput.*, 241:177–196, 2015.
 59. E. R. Wognsen, R. R. Hansen, K. G. Larsen, and P. Koch. Energy-aware scheduling of FIR filter structures using a timed automata model. In *2016 IEEE 19th International Symposium on Design and Diagnostics of Electronic Circuits Systems (DDECS)*, pages 1–6, April 2016.
 60. H. Zhang, C. J. R. Sheppard, T. E. Lipman, and S. J. Moura. Joint fleet sizing and charging system planning for autonomous electric vehicles. *IEEE Transactions on Intelligent Transportation Systems*, 2019.

Technical Appendix

A Proofs

A.1 Proof of Theorem 2

Lemma 8. *There exists a memory-less optimal strategy for the objective MinReach_T .*

Proof. It is clear that in every state s , the player must play a *good* action, i.e. an action a such that $C(s, a) + \max_{s' \in \text{Succ}(s, a)} \text{MinReach}_T(s') \leq \text{MinReach}_T(s)$. If there are multiple good actions in some state s we proceed as follows.

We first assign a *ranking* to states. All states in T have rank 0. Now assume that we have assigned ranks $\leq i$ and we want to assign rank $i + 1$ to some of the yet unranked states. We say that an action a is *progressing* in an unranked state s if all $s' \in \text{Succ}(s, a)$ have a rank (smaller than $i + 1$). If all good actions in the yet unranked states are non-progressing, then all the unranked states s have $\text{MinReach}_T(s) = \infty$, since by playing any good action, the player cannot force reaching a ranked state and thus also a target state; hence, in this case we assign all the unranked states the rank ∞ and finish the construction. Otherwise, we assign rank $i + 1$ to all unranked states that have a good progressing action and continue with the construction. It is easy to see that a state s is assigned a finite rank if and only if $\text{MinReach}_T(s) < \infty$.

We now fix a memory-less strategy σ such that in a state of finite rank, σ chooses a good progressing action; for states of an infinite rank, σ chooses an arbitrary (but fixed) action. Since σ only uses good actions, a straightforward induction shows that for each state s and each s -initiated run that reaches T actually reaches T with consumption at most $\text{MinReach}_T(s)$. So we need to show that σ does not admit runs initiated in a state of finite rank that never reach T . But all σ -compatible runs initiated in a state of finite rank decrease the rank in every step, since σ only plays progressing actions. The result follows. \square

Given a target set T , a number $i \in \mathbb{N}$, and a run ϱ , we define $\text{ReachCons}_T^i(\varrho)$ as $\text{ReachCons}_T(\varrho)$ with the additional restriction that $f \leq i$. Intuitively, $\text{ReachCons}_T^i(\varrho) = \infty$ if ϱ does not visit T within $i - 1$ steps. We then put

$$\text{MinReach}_T^i(s) = \inf_{\sigma} \sup_{\varrho \in \text{Comp}(\sigma, s)} \text{ReachCons}_T^i(\varrho).$$

Lemma 9. *For every $i \geq 0$ it holds that $\text{MinReach}_T^i(s) = \mathcal{F}^i(\mathbf{x}_T)(s)$.*

Proof. By an induction on i . The base case is simple. Now assume that the equality holds for some $i \geq 0$. For $i + 1$ we get:

$$\begin{aligned} \text{MinReach}_T^{i+1}(s) &= \inf_{\sigma} \sup_{\varrho \in \text{Comp}(\sigma, s)} \text{ReachCons}_T^{i+1}(\varrho) \\ &= \min_{a \in A} \left(C(s, a) + \max_{t \in \text{Succ}(s, a)} \inf_{\sigma} \sup_{\varrho \in \text{Comp}(\sigma, t)} \text{ReachCons}_T^i(\varrho) \right) \\ &\stackrel{(\text{I.H.})}{=} \min_{a \in A} \left(C(s, a) + \max_{t \in \text{Succ}(s, a)} \mathcal{F}^i(\mathbf{x}_T)(t) \right) = \mathcal{F}^{i+1}(\mathbf{x}_T)(s) \end{aligned}$$

\square

Theorem 2. Denote by n the length of the longest simple path in \mathcal{M} . Let \mathbf{x}_T be a vector such that $\mathbf{x}_T(s) = 0$ if $s \in T$ and $\mathbf{x}_T(s) = \infty$ otherwise. Then iterating \mathcal{F} on \mathbf{x}_T yields a fixpoint in at most n steps and this fixpoint equals MinReach_T .

Proof. It is easy to see that for each s and each $i \geq 0$ it holds that $\text{MinReach}_T^i(s) \geq \text{MinReach}_T(s)$. Furthermore, an easy computation shows that $\text{MinReach}_T(s)$ is a fixed point of \mathcal{F} . Hence, it suffices to show that $\text{MinReach}_T^{|S|}(s) = \text{MinReach}_T(s)$. Assume, for the sake of contradiction, that we have some $s \in S$ such that $\text{MinReach}_T^{|S|}(s) > \text{MinReach}_T(s)$. Fix σ to be the memory-less optimal strategy from Lemma 8. By the definition of $\text{MinReach}_T^{|S|}(s)$ we have a run $\varrho \in \text{Comp}(\sigma, s)$ such that $\text{ReachCons}_T^{|S|}(\varrho) > \text{MinReach}_T(s)$. This is only possible if $\text{ReachCons}_T^{|S|}(\varrho) = \infty$, otherwise this would contradict the optimality of σ (note that if $\text{ReachCons}_T^i(\varrho) < \infty$, then $\text{ReachCons}_T^j(\varrho) = \text{ReachCons}_T^i(\varrho)$ for all $j \geq i$). Hence, there is a run compatible with σ whose prefix of length $|S|$ does not contain a target state. But this prefix must contain a cycle, and since σ is memory-less, we can iterate this cycle forever. Hence, the optimal strategy σ admits a run from s that never reaches T , a contradiction with $\text{MinReach}_T(s) < \text{MinReach}_T^{|S|}(s) \leq \infty$. \square

A.2 Proof of Lemma 3

By induction on i . The base case is clear. Now assume that the equality holds for some $i \geq 0$. Note that for any $(s, a) \in S \times A$ we have $\text{Succ}(\tilde{s}, a) = \text{Succ}(s, a) \subseteq S$. Also, for all $s' \in S \setminus R$ we have $\llbracket \mathcal{G}^i(\infty) \rrbracket(s') = \mathcal{G}^i(\infty)(s') = \mathcal{F}^i(\mathbf{x}_R)(s')$ (by induction hypothesis), while for $s' \in S \cap R$ we have $\llbracket \mathcal{G}^i(\infty) \rrbracket(s') = 0 = \mathcal{F}^i(\mathbf{x}_R)(s')$ (by the definition of \mathcal{F}). Hence, $\max_{s' \in \text{Succ}(s, a)} \llbracket \mathcal{G}^i(\infty) \rrbracket(s') = \max_{s' \in \text{Succ}(\tilde{s}, a)} \mathcal{F}^i(\mathbf{x}_R)(s')$. Thus,

$$\begin{aligned} \mathcal{G}^{i+1}(\infty)(s) &= \min_{a \in A} \left(C(s, a) + \max_{s' \in \text{Succ}(s, a)} \llbracket \mathcal{G}^i(\infty) \rrbracket(s') \right) \\ &= \min_{a \in A} \left(C(s, a) + \max_{s' \in \text{Succ}(\tilde{s}, a)} \mathcal{F}^i(\mathbf{x}_R)(s') \right) \\ &= \mathcal{F}^{i+1}(\mathbf{x}_R)(\tilde{s}). \end{aligned}$$

(The last equation following from the definition of \mathcal{F} and from the fact that \tilde{s} is never a reload state.)

Moreover, from Lemma 9 it follows that $\mathcal{F}^{i+1}(\mathbf{x}_R)(s) = \text{MinReach}_R^{i+1}(s)$. But if s is not a reload state, then $\text{MinReach}_R^{i+1}(s) = \text{MinReach}_R^{i+1}(\tilde{s})$, so for $s \in S \setminus R$ we have $\mathcal{G}^{i+1}(\infty)(s) = \mathcal{F}^{i+1}(\mathbf{x}_R)(\tilde{s}) = \mathcal{F}^{i+1}(\mathbf{x}_R)(s)$, which proves the second part. \square

A.3 Completion of the proof of Theorem 4

Now we prove that upon termination, $\mathbf{out} \geq \text{Safe}_{\mathcal{M}}$. For every $t \in S$ there exists, by the definition of MinInitCons , a strategy σ_t s.t. $\text{ReachCons}_{\mathcal{M}(\text{Rel}), \text{Rel}}^+(\sigma_t, t) = \mathbf{mic}(t) = \mathbf{out}(t)$ and thus $\text{ReachCons}_{\mathcal{M}(\text{Rel}), \text{Rel}}^+(\varrho)$ is bounded by $\mathbf{out}(t)$ for each $\varrho \in \text{Comp}(\sigma_t, t)$. We construct a new strategy π that, starting in some state t initially mimics σ_t until the next

visit of a state $r_1 \in Rel$. Once this happens, the strategy π begins to mimic σ_{r_1} until the next visit of some $r_2 \in Rel$, when π begins to mimic σ_{r_2} , and so on *ad infinitum*.

Fix any state s such that upon termination, $\mathbf{mic}(s) \leq cap$ (for other states, the inequality $\mathbf{out}(s) \geq Safe_M(s)$ clearly holds). We prove that every run $\varrho \in \text{Comp}_{M(Rel)}(\pi, s)$ is actually $\mathbf{mic}(s)$ -safe from s in the original MDP M . In fact, it is sufficient to show this in $M(Rel)$, since each its reload state is also a reload state of M .

Let $i_1 < i_2 < \dots$ be all indices i such that $\varrho_{i_i} \in Rel$. Since π mimics σ_s until ϱ_{i_1} , we have that $RL_{\mathbf{mic}(s)}(\varrho_{..j}) = \mathbf{mic}(s) - cons(\varrho_{..j})$ where $cons(\varrho_{..j})$ is by definition of σ_s bounded by $\mathbf{mic}(s)$, for all $j \leq i_1$. Now let $m \geq 1$, we set $k = i_m$ and $l = i_{m+1}$. As π mimics σ_{ϱ_k} for between ϱ_k and ϱ_l we have for $k < j \leq l$ that $RL_{\mathbf{mic}(s)}(\varrho_{..j}) = RL_{cap}(\varrho_{k..j}) = cap - cons(\varrho_{k..j})$, where $cons(\varrho_{k..j})$ is bounded by $ReachCons_{Rel}^+(\sigma_{\varrho_k}, \varrho_k) = \mathbf{mic}(\varrho_k) < cap$. Therefore, $RL_{\mathbf{mic}(s)}(\varrho_{..j}) \geq 0$ for all j and ϱ is $\mathbf{mic}(s)$ -safe. This finishes the proof. \square

A.4 Proof of Lemma 4

By induction on i . The base case is clear. Now assume that the statement holds for some $i \geq 0$. Fix any s . Denote by $b = \mathcal{B}^{i+1}(\mathbf{y}_T)(s)$ and $d = SafePR_T^{i+1}(s)$. We show that $b = d$. The equality holds whenever $s \in T$, so in the remainder of the proof we assume that $s \notin T$.

We first prove that $b \geq d$. If $b = \infty$, this is clearly true. Otherwise, let a_{\min} be the action minimizing $SPR\text{-}Val(s, a_{\min}, \mathcal{B}^i(\mathbf{y}_T))$ (which equals b if $s \notin R$) and let $t_{\min} \in Succ(s, a_{\min})$ be the successor used to achieve this value. By induction hypothesis, there exists a strategy σ_1 that is $\mathcal{B}^i(\mathbf{y}_T)(t_{\min})$ -safe in t_{\min} and $\mathbb{P}_{t_{\min}}^{\sigma_1}(Reach_T^i) > 0$, and there also exists a strategy σ_2 that is $Safe(t)$ -safe from t for all $t \in Succ(s, a_{\min}), t \neq t_{\min}$.

Consider now a strategy π which, starting in s , plays a_{\min} . If the outcome of a_{\min} is t_{\min} , π starts to mimic σ_1 , otherwise it starts to mimic σ_2 . We claim that π is b -safe in s and that $\mathbb{P}_s^{\pi}(Reach_T^{i+1}) > 0$ by showing the following two points.

1. There is at least one run $\varrho_T \in \text{Comp}(\pi, s)$ that reaches T in $\leq i + 1$ steps.
2. All runs in $\text{Comp}(\pi, s)$ are b -safe.

We construct ϱ_T easily as $\varrho_T = \alpha \odot \varrho'$ where $\alpha = sa_{\min}t_{\min}$ and ϱ' is the witness that $\mathbb{P}_{t_{\min}}^{\sigma_1}(Reach_T^i) > 0$. Now let $\varrho \in \text{Comp}(\pi, s)$ be a run produced by π from s . Then it has to be of the form $\varrho = sa_{\min}t \odot \varrho_\sigma$ where $\varrho_\sigma \in \text{Comp}(\sigma_1, t)$ if $t = t_{\min}$ or $\varrho_\sigma \in \text{Comp}(\sigma_2, t)$ if $t \neq t_{\min}$; in both cases $Safe(t)$ -safe in t . By definition of $SPR\text{-}Val$ and by induction hypothesis, we have for $s \notin R$ that $b \geq C(s, a_{\min}) + Safe(t)$ and thus $RL_b(sa_{\min}t) \geq Safe(t)$ and thus ϱ is b -safe by Lemma 1. If $s \in R$ and $b = 0$, by similar arguments, as $C(s, a_{\min}) + Safe(t) \leq cap$ (otherwise b would be ∞), we have that $RL_b(sa_{\min}t) \geq Safe(t)$ and thus ϱ is b -safe.

Now we prove that $b \leq d$. This clearly holds if $d = \infty$, so in the remainder of the proof we assume $d \leq cap(M)$. By the definition of d there exists a strategy σ s.t. σ is d -safe in s and $\mathbb{P}_s^{\sigma}(Reach_T^{i+1}(d)) > 0$. Let $a = \sigma(a)$ be the action selected by σ in the first step when starting in s . We denote by τ the strategy such that for all histories α we have $\tau(\alpha) = \sigma(saa)$. For each $t \in Succ(s, a)$ we assign a number d_t defined as $d_t = 0$ if $t \in R$ and $d_t = RL_d(sat)$ otherwise.

We finish the proof by proving these two claims:

1. It holds $SPR\text{-}Val(s, a, \mathcal{B}^i(\mathbf{y}_T)) \leq C(s, a) + \max_{t \in Succ(s, a)} d_t$.
2. If $s \notin R$, then $C(s, a) + \max_{t \in Succ(s, a)} d_t \leq d$.

Let us first see, why these claims are indeed sufficient. From (1.) we get $\mathcal{A}(\mathcal{B}^i(\mathbf{y}_T))(s) \leq C(s, a) + \max_{t \in Succ(s, a)} d_t \leq \text{cap}(\mathcal{M})$ (from the definition of $RL_d(sat)$). If $s \in R$, then it follows that $\llbracket \mathcal{A}(\mathcal{B}^i(\mathbf{y}_T)) \rrbracket(s) = 0 \leq d$. If $s \notin R$, then $\llbracket \mathcal{A}(\mathcal{B}^i(\mathbf{y}_T)) \rrbracket(s) = \mathcal{A}(\mathcal{B}^i(\mathbf{y}_T))(s) \leq C(s, a) + \max_{t \in Succ(s, a)} d_t \leq d$, the first inequality shown above and the second coming from (2.).

So let us start with proving (1.). Note that for each $t \in Succ(s, a)$, τ is necessarily d_t -safe from t ; hence, $\text{Safe}(t) \leq d_t$. Moreover, there exists $q \in Succ(s, a)$ s.t. $\mathbb{P}_q^\tau(\text{Reach}_T^i) > 0$; hence, by induction hypothesis it holds $\mathcal{B}^i(\mathbf{y}_T)(q) \leq d_q$. From this and from the definition of $SPR\text{-}Val$ we get

$$\begin{aligned} SPR\text{-}Val(s, a, \mathcal{B}^i(\mathbf{y}_T)) &\leq C(s, a) + \max\{\mathcal{B}^i(\mathbf{y}_T)(q), \text{Safe}(t) \mid t \in Succ(s, a), t \neq q\} \\ &\leq C(s, a) + \max_{t \in Succ(s, a)} d_t. \end{aligned}$$

To finish, (2.) follows immediately from the definition of d_t and $RL_d(sat)$ as $RL_d(sat)$ is always bounded from above by $d - \text{cap}(s, a)$ for $s \notin R$. \square

A.5 Proof of Lemma 5

By Lemma 4, it suffices to show that $\text{SafePR}_T = \text{SafePR}_T^K$. To this end, fix any state s such that $\text{SafePR}_T(s) < \infty$. For the sake of succinctness, we denote $\text{SafePR}_T(s) = d$. To any strategy π that is d -safe in s , we assign its *index*, which is the infimum of all i such that $\mathbb{P}_s^\pi(\text{Reach}_T^i) > 0$. By assumption that $\text{SafePR}_T(s) < \infty$, there is at least one π with a finite index. Let σ be the d -safe in s strategy with minimal index i : we show that the $i \leq K$.

We proceed by a suitable “strategy surgery” on σ . Let α be a history produced by σ from s of length i whose last state belongs to T . Assume, for the sake of contradiction, that $i > K$. This can only be if at least one of the following conditions hold:

- (a) Some reload state is visited twice on α , i.e. there are $0 \leq j < k \leq i + 1$ such that $\alpha_k = \alpha_j \in R$, or
- (b) some state is visited twice with no intermediate visits to a reload state; i.e., there are $0 \leq j < k \leq i + 1$ such that $\alpha_k = \alpha_j$ and $\alpha_\ell \notin R$ for all $j < \ell < k$.

Indeed, if none of the conditions hold, then the reload states partition α into at most $|R| + 1$ segments, each segment containing non-reload states without repetition. This would imply $i = \text{len}(\alpha) \leq K$.

In both cases (a) and (b) we can arrive at a contradiction using essentially the same argument. Let us illustrate the details on case (a): Consider a strategy π such that for every history of the form $\alpha_{..j} \odot \gamma$ for a suitable γ we have $\pi(\alpha_{..j} \odot \gamma) = \sigma(\alpha_{..k} \odot \gamma)$; on all other histories, π mimics σ . Then π is still d -safe in s . Indeed, the behavior only changed on the suffixes of $\alpha_{..j}$, and for each suitable γ we have $RL_d(\alpha_{..j} \odot \gamma) = RL_d(\alpha_{..k} \odot \gamma)$ due to the fact that $\alpha_j = \alpha_k$ is a reload state. But π has index $i - (k - j) < i$, a contradiction with the choice of σ .

For case (b), the only difference is that now the resource level after $\alpha_{..j} \odot \gamma$ can be higher than the one of $\alpha_{..k} \odot \gamma$ due to the removal of the intermediate non-reloading cycle.

Since we need to show that the energy level never drops below 0, the same argument works. \square